

Project

Charan Basireddy, Bala Aakash Malepati, Mary Sravani Thamtam

2023-04-25

Abstract

Airbnb has disrupted the hospitality industry by providing a platform for short-term rentals. Understanding the factors that determine Airbnb prices is essential for both hosts and guests. This report explores the determinants of Airbnb prices in Europe using data from the publicly available dataset ‘Airbnb Price Determinants’ from Kaggle. In this study, we have explored linear and polynomial regression and random forest models to determine the price of a room and found the main determinants based on the coefficient values.

Introduction

Airbnb has grown rapidly in Europe, providing a popular alternative to traditional hotels and accommodations. The platform enables hosts to rent out their homes or rooms to guests, often at lower prices than hotels. Airbnb offers guests the opportunity to experience local neighborhoods and culture more authentically. For hosts, Airbnb provides a source of income and an opportunity to meet people worldwide. However, Airbnb prices can vary widely, and understanding the factors determining these prices is important for hosts and guests.

How can this information be used: Data can help travelers find accommodation that meets their needs without exceeding budget. Can help hosts set competitive pricing and optimize listings to get more bookings. Help investors evaluate the value of investing in real estate in different European cities based on pricing trends.

Methods

To explore the determinants of Airbnb prices in Europe, we used data on Airbnb listings in ten major European cities (Amsterdam, Athens, Barcelona, Berlin, Budapest, Lisbon, London, Paris, Rome, Vienna). The data consists of various attributes about the hotel. We used regression analysis and random forest to determine the factors that are strongly associated with Airbnb prices in Europe.

Exploratory Data Analysis

We did EDA for the following variables




Figure1

The highest prices in Europe are found in Amsterdam.




Figure2

The prices of entire home are high comparatively

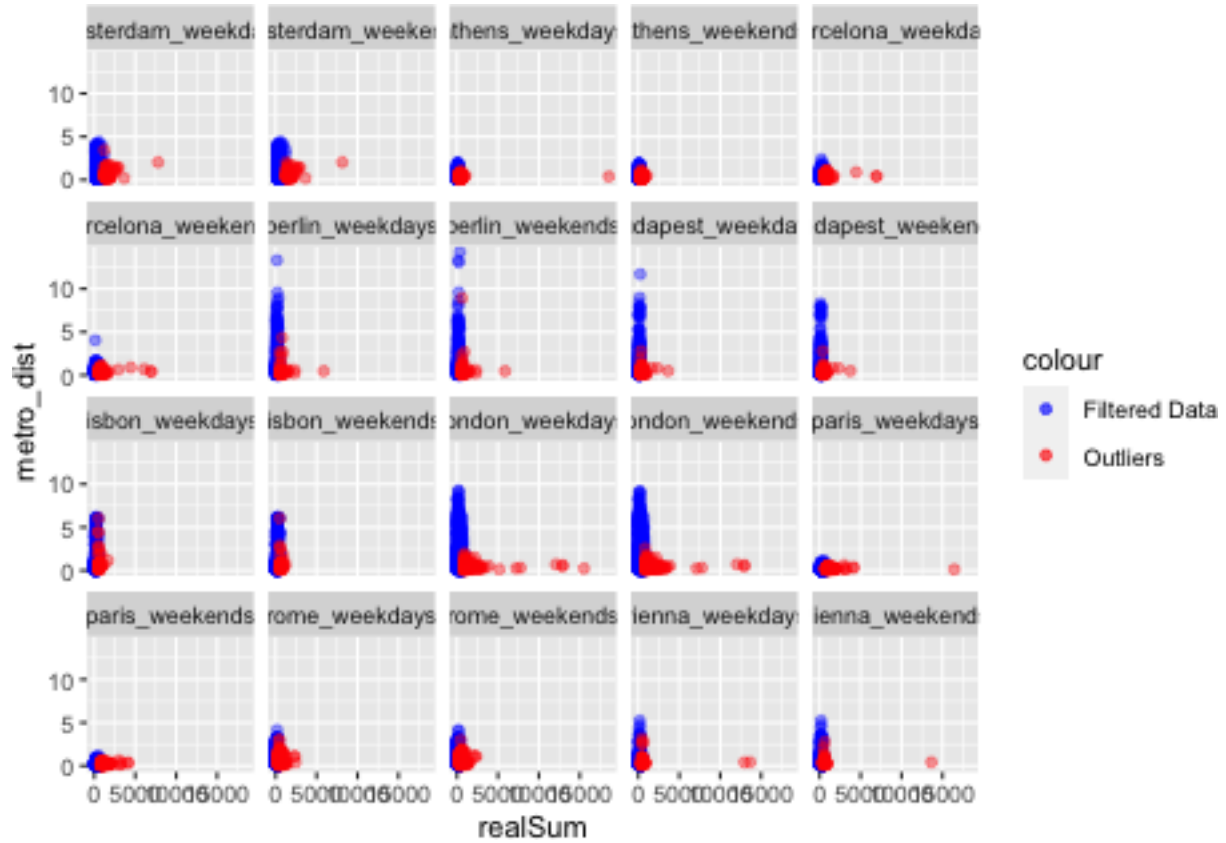


Figure3

In general the rooms that are closer to metro have comparatively higher prices. But, in Rome city the distance to metro is almost same for both categories of price.




Figure4

The attractive index of cities is varying across different cities it is higher for some cities like Paris.




Figure5

There are differences in room prices between both share and non-shared rooms.

Pre Processing

The room_shared and room_private information is already embedded in room_type. The variables are multi-collinear, so we have removed room_shared and room_private.

Multivariate linear regression

Multivariate linear regression is a statistical method that models the relationship between multiple input variables and a continuous output variable. A linear equation is fitted to the input variables, with coefficients representing each input variable's contribution to the output variable. The model assumes that the relationship between the input and output variables is linear and that errors are normally distributed and independent.

Polynomial regression

In polynomial regression, the input variables are raised to various powers. The degree of the polynomial determines the complexity of the model and the number of input variables used in the model. Polynomial

regression is useful when the relationship between the input variables and the output variable is not linear and can provide a better fit to the data than linear regression.

Interaction Variables

Interaction variables, also known as interaction effects, refer to the impact that the combination of two or more input variables has on the output variable in a regression model. They capture the non-additive relationship between input variables and can help to explain better the relationship between the input variables and the output variable. The interaction variables are created by multiplying the values of two or more input variables and including them as additional terms in the model. The coefficient of an interaction variable represents the change in the response variable for a one-unit change in one input variable while holding the other input variables constant.

Random forest regression

Random forest regression is a machine learning technique that combines the power of decision trees with the concept of ensemble learning.

Decision trees work by recursively partitioning the input data into subsets based on the values of the input features to minimize the variance of the target variable within each subgroup. The result is a tree-like structure representing a set of rules for predicting the target variable. Each internal node in the tree corresponds to a test on one of the input features, and each leaf node corresponds to a predicted target variable value. During the prediction phase, the input data is traversed down the tree according to the rules represented by the internal nodes until a leaf node is reached, which provides the predicted value for the target variable. Decision trees handle non-linear relationships between features and the target variable. However, it is prone to overfitting, especially when the tree is deep, and can be sensitive to small changes in the input data.

In Random Forest, multiple decision trees are trained on different subsets of the input data, and the results are combined to make predictions. Each tree in the forest is trained on a random subset of the available features, which helps to reduce overfitting and increase the model's generalization performance. During the prediction phase, the output of each tree is aggregated to produce a final prediction.

Data

Each major city has its own dataset for weekend and weekdays Variables included in data set:

- Host ID (Id)
- Total price of listing (realSum)
- Room type: private, shared, entire home, apt (room_type)
- Whether or not room is shared (room_shared)
- Max number of people allowed in property (person_capacity)
- Whether or not host is superbost (host_is_superhost)
- Whether or not it is multiple rooms (multi)
- Whether for business or family use (biz)
- Distance from city center (dist)
- Distance from nearest metro (metro_dist)
- Latitude and longitude (lat lng)
- Guest satisfaction (guest_satisfaction_overall)
- Cleanliness (cleanliness_rating)
- Total quantity of bedrooms available among all properties for single host (bedrooms)
- Index of Attractions near the hotel (attr_index)

- Normalized Index of Attractions near the hotel (attr_index_norm)
- Index Restaurants near the hotel (rest_index)
- Normalized Index of Restaurants near the hotel (rest_index_norm)

The dataset consists of

- Continuous variables : realSum, dist, metro_dist, lat, lng, attr_index, attr_index_norm, rest_index, rest_index_norm
- Ordinal : person_capacity, guest_satisfaction_overall, cleanliness_rating, bedrooms
- Nominal : room_type, room_shared, host_is_superhost, multi, biz

Results

We have modeled multivariate regression for each city, and below are the coefficients of each model arranged in descending order. The coefficient with a larger value is the essential determinant of the hotel room price.

According to the table, all the models have the same descending order of coefficients. The order is as follows:

1 - room_type 2 - person_capacity 3 - host_is_superhost 4 - multi 5 - biz 6 - cleanliness_rating 7 - guest_satisfaction_overall 8 - bedrooms 9 - dist 10 - metro_dist 11 - attr_index_norm 12 - rest_index_norm, 13 - lng 14 - lat

Modelling

MVLR Seperated by City and Day

Model	var_names	var_coefs	Model	var_names	var_coefs	Model	var_names	var_coefs	Model	var_names	var_coefs	Model	var_names	var_coefs
M_0	(Intercept)	453.805963 M_1	(Intercept)	677.7774655 M_2	(Intercept)	81087.31478 M_3	(Intercept)	3313.326331 M_4	(Intercept)	14214.11864				
M_0	room_typePrivate room	384.5953101 M_1	room_typePrivate room	95.38419868 M_2	room_typePrivate room	148.9417279 M_3	room_typePrivate room	83.28285599 M_4	room_typePrivate room	25.47757592				
M_0	room_typeShared room	136.0046307 M_1	room_typeShared room	42.51664607 M_2	room_typeShared room	78.35566832 M_3	room_typeShared room	32.55688542 M_4	room_typeShared room	12.36521125				
M_0	person_capacity	119.9608563 M_1	person_capacity	31.890407 M_2	person_capacity	38.86525794 M_3	person_capacity	32.36241001 M_4	person_capacity	6.824991537				
M_0	host_is_superhostTrue	42.08786555 M_1	host_is_superhostTrue	16.91133475 M_2	host_is_superhostTrue	18.18910642 M_3	host_is_superhostTrue	20.77860946 M_4	host_is_superhostTrue	5.296527081				
M_0	multi	20.6527087 M_1	multi	4.421357529 M_2	multi	12.07097956 M_3	multi	12.77790391 M_4	multi	1.652498502				
M_0	biz	16.61275571 M_1	biz	2.725004443 M_2	biz	6.985320078 M_3	biz	6.194671034 M_4	biz	0.462524562				
M_0	cleanliness_rating	3.757397212 M_1	cleanliness_rating	2.640302855 M_2	cleanliness_rating	2.492178554 M_3	cleanliness_rating	4.446168901 M_4	cleanliness_rating	-0.492631965				
M_0	guest_satisfaction_overall	2.738196195 M_1	guest_satisfaction_overall	2.335020293 M_2	guest_satisfaction_overall	1.774791262 M_3	guest_satisfaction_overall	1.658915477 M_4	guest_satisfaction_overall	-2.915525964				
M_0	bedrooms	2.025200616 M_1	bedrooms	1.221665363 M_2	bedrooms	0.215433391 M_3	bedrooms	0.831750876 M_4	bedrooms	-13.33265099				
M_0	dist	-2.83720865 M_1	dist	1.053387666 M_2	dist	0.055349136 M_3	dist	0.423821302 M_4	dist	-13.38768559				
M_0	metro_dist	-7.074399447 M_1	metro_dist	-23.4210691 M_2	metro_dist	-9.749578513 M_3	metro_dist	0.346180902 M_4	metro_dist	-17.02253017				
M_0	attr_index_norm	-21.61855664 M_1	attr_index_norm	-24.66793817 M_2	attr_index_norm	-302.2199647 M_3	attr_index_norm	-39.28949375 M_4	attr_index_norm	-47.92355904				
M_0	rest_index_norm	-60.95450514 M_1	rest_index_norm	-29.95296764 M_2	rest_index_norm	-311.0956042 M_3	rest_index_norm	-102.8898274 M_4	rest_index_norm	-48.61561818				
M_0	lng	-175.0198029 M_1	lng	-65.86473148 M_2	lng	-384.8001582 M_3	lng	-120.4611811 M_4	lng	-114.5660309				
M_0	lat	-340.2468974 M_1	lat	-28023.54948 M_2	lat	-1939.911175 M_3	lat	-221.3241991 M_4	lat	-251.7212566				
M_5	(Intercept)	24.49430744 M_6	(Intercept)	6546.714298 M_7	(Intercept)	41416.99566 M_8	(Intercept)	428.4623235 M_9	(Intercept)	408.940146				
M_5	room_typePrivate room	20.76733214 M_6	room_typePrivate room	163.4989324 M_7	room_typePrivate room	130.3445018 M_8	room_typePrivate room	44.47609905 M_9	room_typePrivate room	88.51887853				
M_5	room_typeShared room	8.685270433 M_6	room_typeShared room	36.38305957 M_7	room_typeShared room	81.79791623 M_8	room_typeShared room	12.32210245 M_9	room_typeShared room	36.04103653				
M_5	person_capacity	5.728057336 M_6	person_capacity	25.54307154 M_7	person_capacity	56.22439359 M_8	person_capacity	3.882518636 M_9	person_capacity	20.32350711				
M_5	host_is_superhostTrue	4.819142536 M_6	host_is_superhostTrue	12.82925563 M_7	host_is_superhostTrue	43.15533016 M_8	host_is_superhostTrue	3.699383339 M_9	host_is_superhostTrue	17.90834436				
M_5	multi	1.238674167 M_6	multi	9.982245594 M_7	multi	25.20968259 M_8	multi	3.695177412 M_9	multi	4.390508977				
M_5	biz	0.374695799 M_6	biz	0.933858045 M_7	biz	12.5560253 M_8	biz	1.475050781 M_9	biz	1.60480035				
M_5	cleanliness_rating	0.177646654 M_6	cleanliness_rating	0.207399315 M_7	cleanliness_rating	8.107228556 M_8	cleanliness_rating	0.716854353 M_9	cleanliness_rating	1.352556262				
M_5	guest_satisfaction_overall	0.130448297 M_6	guest_satisfaction_overall	0.113545147 M_7	guest_satisfaction_overall	7.409335312 M_8	guest_satisfaction_overall	0.513052541 M_9	guest_satisfaction_overall	1.242051249				
M_5	bedrooms	-8.337627302 M_6	bedrooms	-19.2800813 M_7	bedrooms	1.641809145 M_8	bedrooms	-0.225693578 M_9	bedrooms	-12.84826556				
M_5	dist	-12.56848529 M_6	dist	-21.77143781 M_7	dist	1.3085809 M_8	dist	-0.783530606 M_9	dist	-13.17443006				
M_5	metro_dist	-59.19710383 M_6	metro_dist	-22.96385155 M_7	metro_dist	-4.603578145 M_8	metro_dist	-1.508888739 M_9	metro_dist	-13.85744129				
M_5	attr_index_norm	-82.58885424 M_6	attr_index_norm	-130.3824975 M_7	attr_index_norm	-97.88170929 M_8	attr_index_norm	-6.252414347 M_9	attr_index_norm	-16.13640861				
M_5	rest_index_norm	-190.41590079 M_6	rest_index_norm	-181.4255816 M_7	rest_index_norm	-284.4850304 M_8	rest_index_norm	-39.8036185 M_9	rest_index_norm	-28.48935415				
M_5	lng	-474.7694686 M_6	lng	-193.4634265 M_7	lng	-715.552286 M_8	lng	-94.72375311 M_9	lng	-120.2455734				
M_5	lat	-2021.541971 M_6	lat	-212.9519776 M_7	lat	-821.2043548 M_8	lat	-17868.39131 M_9	lat	-20955.97861				

M_0 - Amsterdam, M_1 - Athens, M_2 - Barcelona, M_3 - Berlin, M_4 - Budapest, M_5 - Lisbon, M_6 - London, M_7 - Paris, M_8 - Rome, M_9 - Vienna

MVLR Combined of all Cities

term	estimate	std.error	statistic	p.value
city_dayathens_weekdays	6315.09061077375	1388.53511239174	4.54802370816251	5.43276632399387E-06
city_dayathens_weekends	6303.53106566887	1388.65271393214	4.53931425937278	5.66191686814906E-06
city_daybudapest_weekends	3929.67339172125	706.843976888775	5.55946364432218	2.72512019627685E-08
city_daybudapest_weekdays	3902.35110769949	706.88057311247	5.52052391327863	3.40308057888361E-08
city_dayvienna_weekdays	3231.8185121282	582.709191916445	5.54619449454577	2.93992522226129E-08
city_dayvienna_weekends	3230.26798535313	582.79716507451	5.54269680591213	2.9992285987412E-08
city_dayrome_weekends	2952.53521434873	881.429677236262	3.34971159991624	0.000809784171597337
city_dayrome_weekdays	2947.68523500212	881.398370960292	3.34432798167143	0.00082565958973142
city_dayberlin_weekends	1958.88440687424	342.040071097647	5.72706116154153	1.02993861536248E-08
city_dayberlin_weekdays	1949.04242701527	342.124495503847	5.69688067539542	1.22965025210763E-08
city_daybarcelona_weekends	429.652941610428	837.810922281234	0.512828050081453	0.608074739358428
city_daybarcelona_weekdays	411.771664434536	837.79087513094	0.491496955454641	0.623077988219532
lat	123.211714028539	76.5227515358753	1.61013177853093	0.107377819824515
bedrooms	86.0154191816913	3.18878180193731	26.9743822325609	1.10487900248437E-158
city_dayamsterdam_weekends	67.9410312574828	16.0017245274658	4.24585682254852	2.18302692822801E-05
biz	33.2805902980968	4.18846537329009	7.94577185962373	1.98543351601859E-15
person_capacity	23.962639711116	1.76448699475381	13.580513646381	6.6229089695509E-42
multi	9.60044781984104	4.13239391830793	2.32321700438764	0.0201730153901032
attr_index_norm	6.37045498927847	0.294558501285674	21.6271299639054	4.50013019203766E-103
cleanliness_rating	5.03832622544896	2.41532715274561	2.08598086587221	0.0369873477630624
host_is_superhostTrue	1.07487304905057	3.93436342704325	0.273201260885642	0.784700071060939
guest_satisfaction_overall	0.776010513381954	0.261452270487271	2.96807716351323	0.00299865383859512
rest_index_norm	-0.183747213632232	0.177364361505566	-1.03598723031214	0.300215028045804
dist	-1.53304759323782	1.26282244844453	-1.21398506585478	0.224761355845419
metro_dist	-3.99666546744126	2.50252382009665	-1.59705391626878	0.110262427719338
room_typePrivate room	-114.36546946963	4.28228402903422	-26.7066520329393	1.28065674381786E-155
room_typeShared room	-204.184154195175	18.9348057155667	-10.7835357416586	4.52732676348359E-27
lng	-262.890929428664	40.1930553101859	-6.54070528850891	6.20447859714726E-11
city_dayparis_weekdays	-403.028852808039	278.481924100759	-1.44723523478032	0.147839720417429
city_dayparis_weekends	-422.138907209481	278.643708051156	-1.51497735284222	0.129786880370265
city_daylondon_weekdays	-1409.29965386078	206.104550214087	-6.83779010408502	8.17001208243907E-12
city_daylondon_weekends	-1410.93275121758	206.122303369817	-6.84512412364295	7.76270471372057E-12
city_daylisbon_weekends	-2304.00669656451	1143.81264744771	-2.0143217525227	0.0439831538834079
city_daylisbon_weekdays	-2312.9309439741	1143.89771522079	-2.02197356738986	0.0431864319357264
(Intercept)	-4954.90756116218	3996.18236913533	-1.2399102702198	0.215016631100909

Apart from Cityday, below are the ranking of coefficients of the model

1. Lat,
2. bedrooms,
3. biz,
4. person_capacity,
5. Multi,
6. Attr_index_norm,
7. Cleanliness_rating,

8. host_is_superhost, 9. guest_satisfaction_overall, 10. rest_index_norm 11. dist, 12. metro_dist, 13. room_type, 14. lng

The city day coefficients are almost ranked at the top.

Other Models


Model	Train R^2	Train Adj R^2	Train $RMSE$	Test R^2	Test Adj R^2	Test $RMSE$
MLR	0.2150414	0.2146725	304.9175	0.33744	0.3371287	233.2618
MLR with Step	0.2149964	0.2146275	304.9262	0.3373441	0.3370327	233.2787
MLR with IVs	0.2159815	0.215613	304.7348	0.3379444	0.3376333	233.173
MLR with IVs, Step	0.2159435	0.215575	304.7422	0.3380392	0.3377281	233.1563
Poly with Order 2	0.2154699	0.2151012	304.8342	0.3373538	0.3370424	233.277
Poly with Order 2, Step	0.2154289	0.2150602	304.8422	0.3373204	0.337009	233.2829
Poly with Order 2, IVs	0.22214	0.2217744	303.5356	0.334993	0.3346805	233.6922
Poly with Order 2, IVs, Step	0.2221344	0.2217689	303.5367	0.3350694	0.3347569	233.6788
Poly with Order 3	0.2160624	0.215694	304.7191	0.3375855	0.3372742	233.2362
Poly with Order 3, Step	0.2160337	0.2156653	304.7247	0.3376519	0.3373407	233.2245
Poly with Order 3, IVs	0.2330663	0.2327059	301.3962	0.1901115	0.189731	257.8955
Poly with Order 3, IVs, Step	0.2285384	0.2281759	302.2846	0.3281442	0.3281442	234.8925
Lasso With Order 1	0.215693	0.2147166	325.7776	0.3372306	0.3353023	233.2987
Lasso With Order 2	0.2201963	0.217819	326.6683	0.1830879	0.1772535	259.0113
Lasso With Order 3	0.2228371	0.2166688	326.2303	0.02854527	0.01036279	282.4505
Random Forest	0.8747755	0.8724317	121.7878	0.7588743	0.7543612	140.719

* IVs here is Interaction variables

Random Forest

We have trained Random Forest on all cities' combined data, and below are the important variables.

	IncNodePurity
room_type	122087850
host_is_superhost	26594131
multi	21575815
biz	24190212
city_day	114299029
person_capacity	183757090
cleanliness_rating	79985960
guest_satisfaction_overall	162106050
bedrooms	342017315
dist	423128244
metro_dist	248206656
attr_index_norm	662153431
rest_index_norm	435838043
lng	482308349
lat	588248566



The random forest model has chosen the attraction index as the most important variable, which makes more sense from a general point of view because the places with more attractions would attribute to pricing differences. The other variables ranking also seems consistent with common sense, such as latitude and longitude, which indicates location and distance from the city center, etc.

Overall the random forest has performed better than the other models with a Adjusted. R^2 value of 0.7543 and RMSE of 140.719 compared to the best regression model, which has Adjusted R^2 value of 0.3377 and RMSE of 233.1563 on the test set.

Discussion

When we applied separate models for each city data, on linear regression, latitude was least ranked, but when we combined all the data latitude is on top rank next to city day, it indicates the information about location is being captured by the latitude when all data is combined, in separate models all the data about city represents its location so latitude is least ranked.

The attraction index is ranked higher in all cities combined model than separated models. (From *Figure 4*)

The host_is_super host is ranked high in individual models than combined model, but its significance value is very low.

Room Type is ranked at the top in separate models, but in the combined model is at quite a low rank. It is due to the information of Room Type capturing the information about the location; it can be inferred from the graph *Figure 5*

Summary

This report investigates the factors that determine Airbnb room prices in Europe using data from the publicly available dataset ‘Airbnb Price Determinants.’ To identify the main determinants of Airbnb room prices in Europe, the report employed a regression analysis and random forest model. The dataset was collected from ten major European cities, including Amsterdam, Athens, Barcelona, Berlin, Budapest, Lisbon, London, Paris, Rome, and Vienna. Exploratory data analysis was conducted on various attributes of the data. The report discovered that Amsterdam has the highest Airbnb room prices in Europe. In general, entire homes have higher prices compared to shared rooms. The proximity of a room to a metro station was found to be correlated with higher prices. However, in Rome, the distance to the metro station has a negligible impact on room prices. The study also discovered that there are differences in room prices between shared and non-shared rooms. Furthermore, the study used multivariate linear regression, polynomial regression, interaction variables, and random forest regression to identify the factors that are strongly associated with Airbnb prices in Europe. The report demonstrates that the main determinants of Airbnb room prices in Europe include the number of people allowed in a room, the room type, and the host’s superhost status. The report provides valuable insights for Airbnb guests, hosts, and investors in evaluating the value of investing in real estate in different European cities based on pricing trends.

Appendix for Code and Detailed Analysis

Pre Processing and Cleaning the Data

Data loading

```
# Set the relative directory path
my_dir <- "../archive"

# List all the files in the directory
files <- list.files(path = my_dir, full.names = TRUE)
```

Combining the Data from all Files

```

# Get a list of all the csv files in the directory
file_list <- list.files(path = my_dir, pattern = "*.csv", full.names = TRUE)

# Initialize an empty list to store the data frames
df_list <- list()

# Loop through each file and read it into a data frame
for (i in seq_along(file_list)) {
  df <- read.csv(file_list[i])

  # Add a new column with the city_day
  df$city_day <- basename(file_list[i])

  # Append the data frame to the list
  df_list[[i]] <- df
}

# Combine all the data frames into a single dataset
my_data <- bind_rows(df_list)

# Removing the .csv ext
my_data$city_day <- gsub("\\.csv", "", my_data$city_day)

# Print the first few rows of the data
head(my_data)

```

```

print(unique(my_data[my_data$room_shared == my_data$room_private,
  ]$room_type)) # if the room is shared

```

```
## [1] "Entire home/apt"
```

```
print(unique(my_data[my_data$room_private == "False", ]$room_type))
```

```
## [1] "Entire home/apt" "Shared room"
```

```
print(unique(my_data[my_data$room_shared == "True", ]$room_type))
```

```
## [1] "Shared room"
```

```
print(unique(my_data[my_data$room_shared == "False", ]$room_type))
```

```
## [1] "Private room" "Entire home/apt"
```


The `room_shared` and `room_private` information is already embedded in `room_type`. The variables are multi-collinear, so we can remove `room_shared` and `room_private`.

Dropping columns of `room_shared` and `room_private`


```
my_data = select(my_data, -c(room_shared, room_private))
head(my_data)
```

```
##   X   realSum   room_type person_capacity host_is_superhost multi biz
## 1 0 194.0337 Private room           2           False      1   0
## 2 1 344.2458 Private room           4           False      0   0
## 3 2 264.1014 Private room           2           False      0   1
## 4 3 433.5294 Private room           4           False      0   1
## 5 4 485.5529 Private room           2            True      0   0
## 6 5 552.8086 Private room           3           False      0   0
##   cleanliness_rating guest_satisfaction_overall bedrooms      dist metro_dist
## 1              10              93           1 5.0229638 2.5393800
## 2              8              85           1 0.4883893 0.2394039
## 3              9              87           1 5.7483119 3.6516213
## 4              9              90           2 0.3848620 0.4398761
## 5             10              98           1 0.5447382 0.3186926
## 6              8             100           2 2.1314201 1.9046682
##   attr_index attr_index_norm rest_index rest_index_norm      lng      lat
## 1   78.69038     4.166708   98.25390     6.846473 4.90569 52.41772
## 2  631.17638    33.421209  837.28076    58.342928 4.90005 52.37432
## 3   75.27588     3.985908   95.38695     6.646700 4.97512 52.36103
## 4  493.27253    26.119108  875.03310    60.973565 4.89417 52.37663
## 5  552.83032    29.272733  815.30574    56.811677 4.90051 52.37508
## 6  174.78896     9.255191  225.20166    15.692376 4.87699 52.38966
##           city_day
## 1 amsterdam_weekdays
## 2 amsterdam_weekdays
## 3 amsterdam_weekdays
## 4 amsterdam_weekdays
## 5 amsterdam_weekdays
## 6 amsterdam_weekdays
```

```
ggplot() + geom_point(data = my_data, aes(x = attr_index, y = attr_index_norm),
  alpha = 0.4)
```




```
ggplot() + geom_point(data = my_data, aes(x = attr_index, y = attr_index_norm),  
  alpha = 0.4) + facet_wrap(~city_day)
```




attr_index and attr_index_norm are same, attr_index_norm is just normalized attr_index

```
ggplot() + geom_point(data = my_data, aes(x = rest_index, y = rest_index_norm),
  alpha = 0.4)
```

```
ggplot() + geom_point(data = my_data, aes(x = rest_index, y = rest_index_norm),  
  alpha = 0.4) + facet_wrap(~city_day)
```



rest_index and rest_index_norm are same, rest_index_norm is just normalized rest_index.

removing attr_index and rest_index

```
my_data = select(my_data, -c(attr_index, rest_index))
head(my_data)
```

Outliers using IQR Range

Filtering out the Outliers from Data Out of IQR Ranges

```
# Initialize an empty list to store the outliers
outliers_list <- list()

# Initialize an empty list to store the filtered data
# frames
df_list_filtered <- list()

# Loop through each file and read it into a data frame
# after removing outliers
for (i in seq_along(file_list)) {
  df_filtered <- read.csv(file_list[i])

  # Add a new column with the city_day
  df_filtered$city_day <- gsub("\\.csv", "", basename(file_list[i]))
}
```

```

iqr_var1 <- IQR(df_filtered$realSum)

# Calculate the upper and lower bounds for each
# variable
upper_var1 <- quantile(df_filtered$realSum, 0.75) + 1.5 *
  iqr_var1
lower_var1 <- quantile(df_filtered$realSum, 0.25) - 1.5 *
  iqr_var1

# Filter the data based on the upper and lower bounds
# for each variable
filtered_data <- filter(df_filtered, realSum > lower_var1 &
  realSum < upper_var1)

# Append the filtered data frame to the list
df_list_filtered[[i]] <- filtered_data

# Get the rows that were removed while filtering
outliers <- anti_join(df_filtered, filtered_data)

# Append the outliers to the list
outliers_list[[i]] <- outliers
}

# Combine all the filtered data frames into a single
# dataset
my_data_filtered <- bind_rows(df_list_filtered)

# Removing the .csv ext
my_data_filtered$city_day <- gsub("\\.csv", "", my_data_filtered$city_day)

summary(my_data_filtered)

```

```

##           X           realSum      room_type      room_shared
## Min.      : 0      Min.      : 34.78      Length:48970      Length:48970
## 1st Qu.: 645      1st Qu.: 145.23      Class :character      Class :character
## Median :1340      Median : 204.27      Mode  :character      Mode  :character
## Mean      :1621      Mean      : 244.35
## 3rd Qu.:2385      3rd Qu.: 295.27
## Max.      :5378      Max.      :1229.11
## room_private      person_capacity      host_is_superhost      multi
## Length:48970      Min.      :2.00      Length:48970      Min.      :0.0000
## Class :character      1st Qu.:2.00      Class :character      1st Qu.:0.0000
## Mode  :character      Median :3.00      Mode  :character      Median :0.0000
##                               Mean      :3.08                               Mean      :0.2953
##                               3rd Qu.:4.00                               3rd Qu.:1.0000
##                               Max.      :6.00                               Max.      :1.0000
##           biz           cleanliness_rating      guest_satisfaction_overall      bedrooms
## Min.      :0.000      Min.      : 2.000      Min.      : 20.00      Min.      : 0.000
## 1st Qu.:0.000      1st Qu.: 9.000      1st Qu.: 90.00      1st Qu.: 1.000
## Median :0.000      Median :10.000      Median : 95.00      Median : 1.000
## Mean      :0.342      Mean      : 9.384      Mean      : 92.57      Mean      : 1.118

```

```
## 3rd Qu.:1.000 3rd Qu.:10.000 3rd Qu.: 98.00 3rd Qu.: 1.000
## Max. :1.000 Max. :10.000 Max. :100.00 Max. :10.000
## dist metro_dist attr_index attr_index_norm
## Min. : 0.01506 Min. : 0.002301 Min. : 15.15 Min. : 0.9263
## 1st Qu.: 1.48598 1st Qu.: 0.250718 1st Qu.: 133.75 1st Qu.: 6.2341
## Median : 2.66962 Median : 0.416955 Median : 228.54 Median : 11.1929
## Mean : 3.24072 Mean : 0.691774 Mean : 285.15 Mean : 13.0064
## 3rd Qu.: 4.31533 3rd Qu.: 0.749700 3rd Qu.: 374.37 3rd Qu.: 16.9444
## Max. :25.28456 Max. :14.273577 Max. :4513.56 Max. :100.0000
## rest_index rest_index_norm lng lat
## Min. : 19.58 Min. : 0.5928 Min. : -9.22634 Min. : 37.95
## 1st Qu.: 245.42 1st Qu.: 8.5601 1st Qu.: -0.07277 1st Qu.: 41.40
## Median : 512.42 Median : 17.1799 Median : 4.87234 Median : 47.51
## Mean : 611.32 Mean : 22.2861 Mean : 7.40027 Mean : 45.66
## 3rd Qu.: 818.44 3rd Qu.: 32.0321 3rd Qu.: 13.52350 3rd Qu.: 51.47
## Max. :6696.16 Max. :100.0000 Max. :23.78602 Max. :52.64
## city_day
## Length:48970
## Class :character
## Mode :character
##
##
##
```

```
# Combine all the outliers into a single dataset
my_outliers <- bind_rows(outliers_list)

# Removing the .csv ext
my_outliers$city_day <- gsub("\\.csv", "", my_outliers$city_day)

summary(my_outliers)
```

```
## X realSum room_type room_shared
## Min. : 0 Min. : 279.4 Length:2737 Length:2737
## 1st Qu.: 666 1st Qu.: 469.2 Class :character Class :character
## Median :1237 Median : 691.9 Mode :character Mode :character
## Mean :1614 Mean : 915.5
## 3rd Qu.:2310 3rd Qu.: 996.3
## Max. :5374 Max. :18545.5
## room_private person_capacity host_is_superhost multi
## Length:2737 Min. :2.000 Length:2737 Min. :0.000
## Class :character 1st Qu.:4.000 Class :character 1st Qu.:0.000
## Mode :character Median :5.000 Mode :character Median :0.000
## Mean :4.628 Mean :0.221
## 3rd Qu.:6.000 3rd Qu.:0.000
## Max. :6.000 Max. :1.000
## biz cleanliness_rating guest_satisfaction_overall bedrooms
## Min. :0.0000 Min. : 2.000 Min. : 20.00 Min. :0.000
## 1st Qu.:0.0000 1st Qu.: 9.000 1st Qu.: 91.00 1st Qu.:1.000
## Median :0.0000 Median :10.000 Median : 97.00 Median :2.000
## Mean :0.4965 Mean : 9.509 Mean : 93.65 Mean :1.886
## 3rd Qu.:1.0000 3rd Qu.:10.000 3rd Qu.:100.00 3rd Qu.:2.000
## Max. :1.0000 Max. :10.000 Max. :100.00 Max. :6.000
## dist metro_dist attr_index attr_index_norm
```

```
## Min. : 0.01504 Min. :0.006171 Min. : 20.5 Min. : 1.468
## 1st Qu.: 1.04119 1st Qu.:0.218081 1st Qu.: 225.1 1st Qu.: 11.719
## Median : 1.89579 Median :0.352339 Median : 385.0 Median : 17.958
## Mean : 2.30674 Mean :0.498426 Mean : 456.2 Mean : 20.892
## 3rd Qu.: 3.00820 3rd Qu.:0.576430 3rd Qu.: 610.6 3rd Qu.: 25.953
## Max. :21.29515 Max. :8.918036 Max. :2040.4 Max. :100.000
## rest_index rest_index_norm lng lat
## Min. : 27.9 Min. : 0.667 Min. : -9.22476 Min. :37.96
## 1st Qu.: 408.5 1st Qu.: 14.187 1st Qu.: -0.06677 1st Qu.:41.41
## Median : 739.9 Median : 30.001 Median : 4.88384 Median :47.51
## Mean : 904.9 Mean : 31.734 Mean : 7.88764 Mean :45.93
## 3rd Qu.:1269.7 3rd Qu.: 45.426 3rd Qu.:13.44666 3rd Qu.:51.50
## Max. :4183.1 Max. :100.000 Max. :23.75400 Max. :52.58
## city_day
## Length:2737
## Class :character
## Mode :character
##
##
##
```

Percentage of Outliers outside of IQR range.

```
# Create empty table
outliers_table <- data.frame(City_day = character(), Data_Length = numeric(),
  Percent_Outliers = numeric(), stringsAsFactors = FALSE)

# Loop through city_data and fill in table
for (city_day in unique(my_data$city_day)) {
  x = my_data[my_data$city_day == city_day, ]$realSum
  q1 <- quantile(x, 0.25)
  q3 <- quantile(x, 0.75)
  iqr <- IQR(x)
  upper_bound <- q3 + 1.5 * iqr
  lower_bound <- q1 - 1.5 * iqr
  x_no_outliers <- x[x >= lower_bound & x <= upper_bound]
  percent_outliers <- ((length(x) - length(x_no_outliers))/length(x)) *
    100

  # Add row to table
  outliers_table <- rbind(outliers_table, data.frame(City_day = city_day,
    Data_Length = length(x), Percent_Outliers = percent_outliers))
}

# Format table using kable
kable(outliers_table, format = "markdown")
```

City_day	Data_Length	Percent_Outliers
amsterdam_weekdays	1103	5.077063
amsterdam_weekends	977	5.629478
athens_weekdays	2653	5.767056

City_day	Data_Length	Percent_Outliers
athens_weekends	2627	5.405405
barcelona_weekdays	1555	7.524116
barcelona_weekends	1278	8.059468
berlin_weekdays	1284	6.308411
berlin_weekends	1200	6.166667
budapest_weekdays	2074	5.930569
budapest_weekends	1948	5.544148
lisbon_weekdays	2857	3.360168
lisbon_weekends	2906	3.475568
london_weekdays	4614	5.353273
london_weekends	5379	5.521472
paris_weekdays	3130	6.134185
paris_weekends	3558	5.368184
rome_weekdays	4492	5.031167
rome_weekends	4535	5.005513
vienna_weekdays	1738	4.257767
vienna_weekends	1799	4.113396

Spilt Training and Testing Data

```
set.seed(123456789)
my_data_train <- my_data[sample(nrow(my_data), 0.7 * nrow(my_data)),
]
my_data_test <- my_data[setdiff(1:nrow(my_data), rownames(my_data_train)),
]
dim(my_data_train)
```

```
## [1] 36194    17
```

```
dim(my_data_test)
```

```
## [1] 15513    17
```


Exploratory Data Analysis

Outlier Analysis

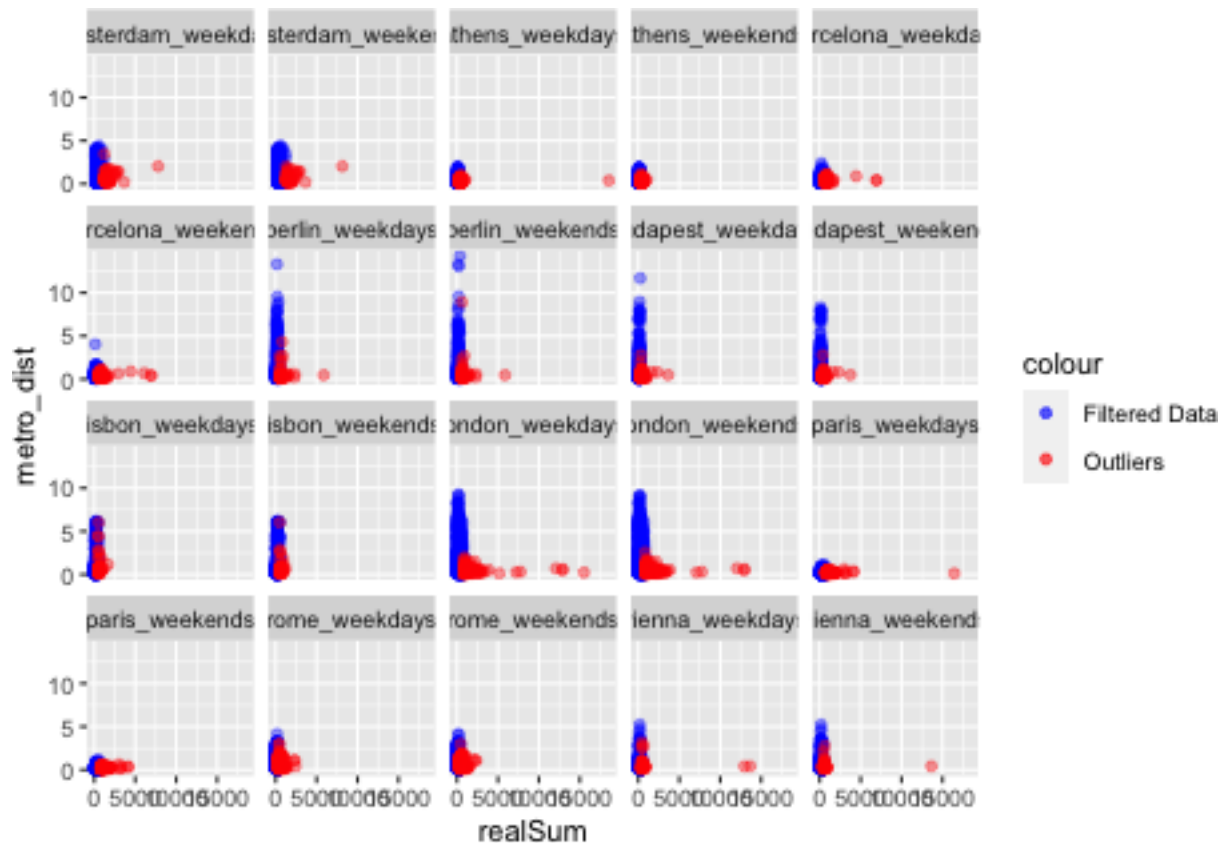
Metro Dist vs Real Sum

We have planned to analyse the filtered data along with outlier data. Here outlier data represents the hotel rooms with high prices.

```
ggplot() + geom_point(data = my_data_filtered, aes(x = realSum,
y = metro_dist, color = "Filtered Data"), alpha = 0.4) +
  geom_point(data = my_outliers, aes(x = realSum, y = metro_dist,
color = "Outliers"), alpha = 0.4) + scale_color_manual(values = c(`Filtered Data` = "blue",
Outliers = "red"))
```




```
ggplot() + geom_point(data = my_data_filtered, aes(x = realSum,
  y = metro_dist, color = "Filtered Data"), alpha = 0.4) +
  geom_point(data = my_outliers, aes(x = realSum, y = metro_dist,
    color = "Outliers"), alpha = 0.4) + scale_color_manual(values = c(`Filtered Data` = "blue",
    Outliers = "red")) + facet_wrap(~city_day)
```



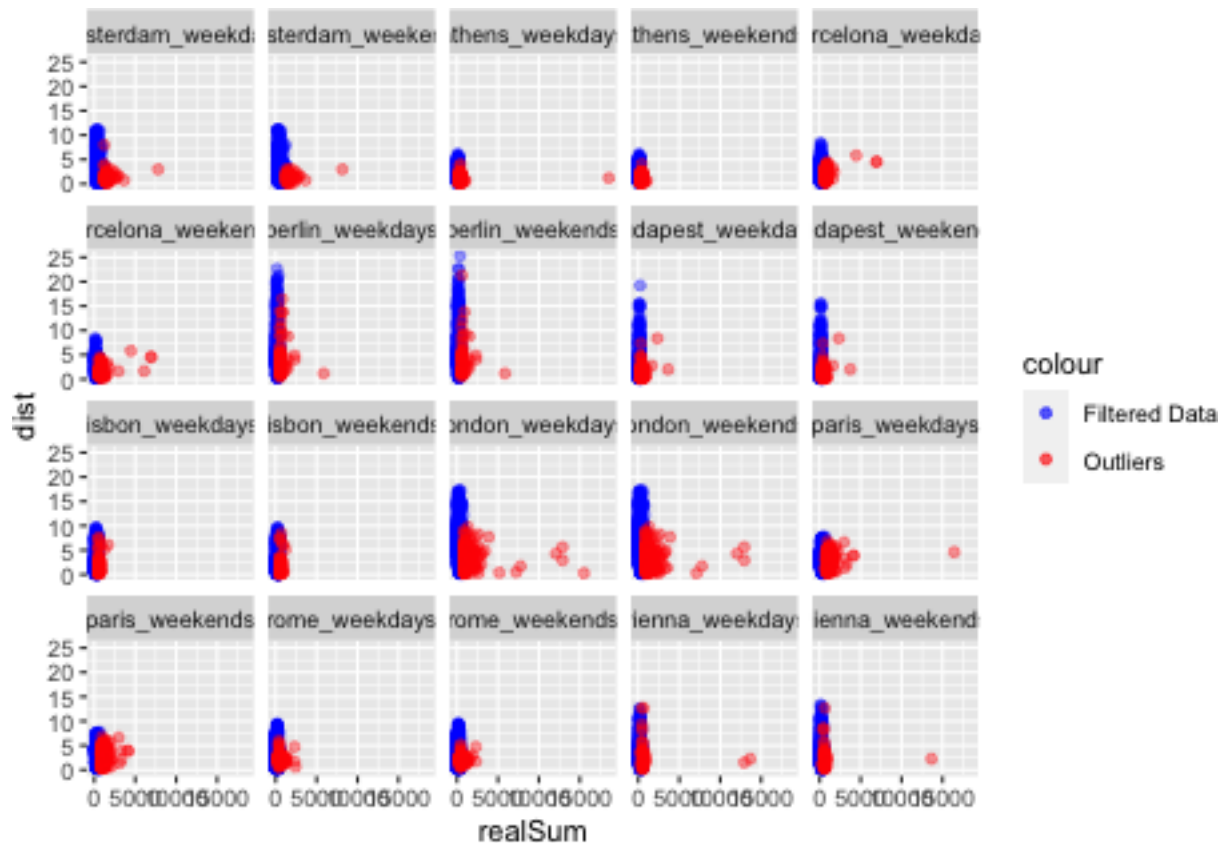
In general the rooms that are closer to metro have comparatively higher prices. But, in Rome city the distance to metro is almost same for both categories of price.

Real Sum vs Distance

```
ggplot() + geom_point(data = my_data_filtered, aes(x = realSum,
  y = dist, color = "Filtered Data"), alpha = 0.4) + geom_point(data = my_outliers,
  aes(x = realSum, y = dist, color = "Outliers"), alpha = 0.4) +
  scale_color_manual(values = c(`Filtered Data` = "blue", Outliers = "red"))
```


```
ggplot() + geom_point(data = my_data_filtered, aes(x = realSum,
  y = dist, color = "Filtered Data"), alpha = 0.4) + geom_point(data = my_outliers,
  aes(x = realSum, y = dist, color = "Outliers"), alpha = 0.4) +
  scale_color_manual(values = c(`Filtered Data` = "blue", Outliers = "red")) +
  facet_wrap(~city_day)
```




In general the pricey rooms are near to the centre of the city.

Real Sum vs Attraction Index Normal


```
ggplot() + geom_point(data = my_data_filtered, aes(x = realSum,
  y = attr_index_norm, color = "Filtered Data"), alpha = 0.4) +
  geom_point(data = my_outliers, aes(x = realSum, y = attr_index_norm,
    color = "Outliers"), alpha = 0.4) + scale_color_manual(values = c(`Filtered Data` = "blue",
    Outliers = "red"))
```



```
ggplot() + geom_point(data = my_data_filtered, aes(x = realSum,
  y = attr_index_norm, color = "Filtered Data"), alpha = 0.4) +
  geom_point(data = my_outliers, aes(x = realSum, y = attr_index_norm,
    color = "Outliers"), alpha = 0.4) + scale_color_manual(values = c(`Filtered Data` = "blue",
  Outliers = "red")) + facet_wrap(~city_day)
```




```
ggplot() + geom_point(data = my_data_filtered, aes(x = realSum,
y = attr_index_norm), alpha = 0.4) + geom_point(data = my_outliers,
aes(x = realSum, y = attr_index_norm), alpha = 0.4)
```



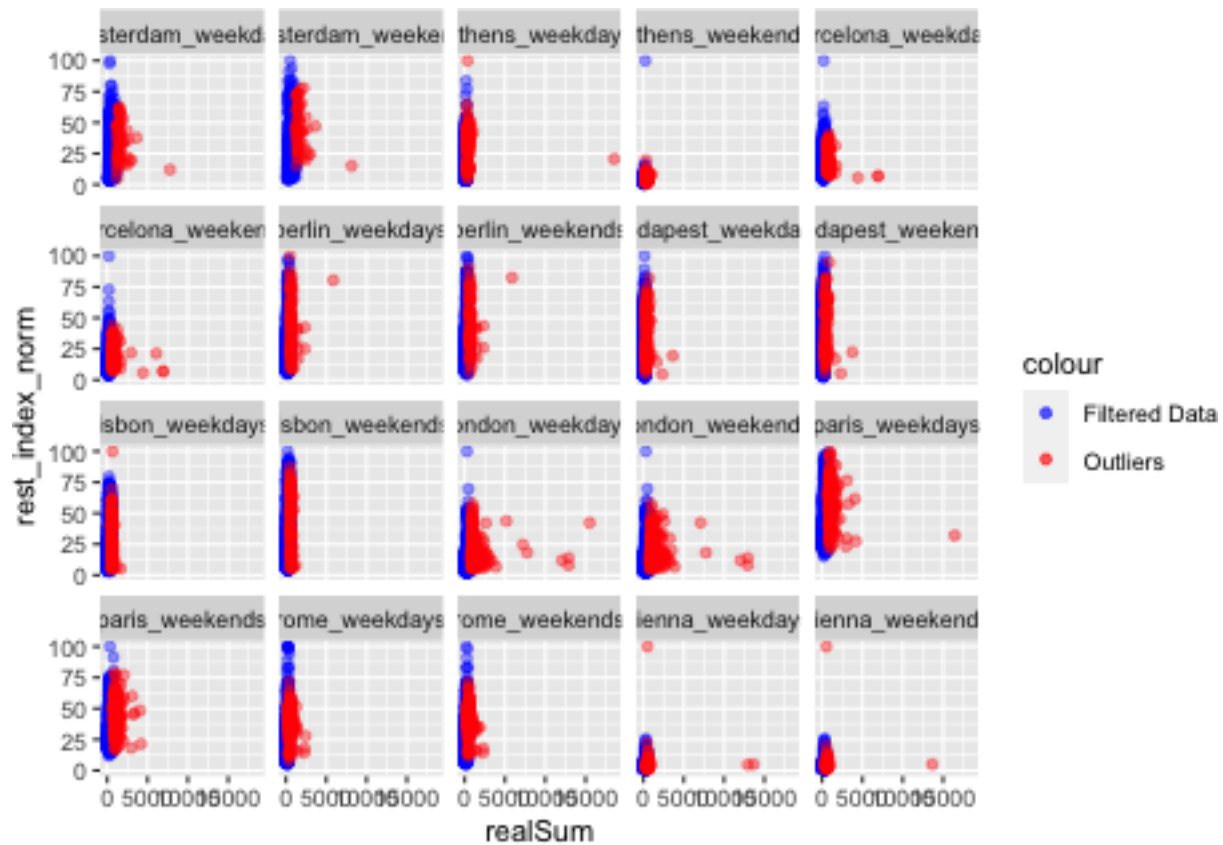
The range of values falling b/w outliers and normal data is almost same . So there isn't a relationship b/w attr_index and realSum.

Real Sum vs Restaurant Index Normal

```
ggplot() + geom_point(data = my_data_filtered, aes(x = realSum,
  y = rest_index_norm, color = "Filtered Data"), alpha = 0.4) +
  geom_point(data = my_outliers, aes(x = realSum, y = rest_index_norm,
    color = "Outliers"), alpha = 0.4) + scale_color_manual(values = c(`Filtered Data` = "blue",
    Outliers = "red"))
```




```
ggplot() + geom_point(data = my_data_filtered, aes(x = realSum,
  y = rest_index_norm, color = "Filtered Data"), alpha = 0.4) +
  geom_point(data = my_outliers, aes(x = realSum, y = rest_index_norm,
    color = "Outliers"), alpha = 0.4) + scale_color_manual(values = c(`Filtered Data` = "blue",
    Outliers = "red")) + facet_wrap(~city_day)
```



There is no relationship between outliers and rest_index

Room Type Vs Real Sum


```
ggplot(my_data, aes(x = realSum, fill = room_type, group = room_type)) +
  geom_histogram(alpha = 0.5, nbins = 20) + theme(axis.title.x = element_text(size = 14),
  axis.title.y = element_text(size = 14))
```



```
ggplot(my_data_filtered, aes(x = realSum, fill = room_type, group = room_type)) +
  geom_histogram(alpha = 0.5, nbins = 20) + theme(axis.title.x = element_text(size = 14),
axis.title.y = element_text(size = 14))
```




```
ggplot(my_data_filtered, aes(x = realSum, fill = room_type, group = room_type)) +
  geom_histogram(alpha = 0.5, nbins = 20) + theme(axis.title.x = element_text(size = 14),
axis.title.y = element_text(size = 14)) + facet_wrap(~city_day)
```


The price of entire home/apt tend to be higher compared to other two categories. And the count of entire home /apt is also more.

Room Type Vs Person Capacity

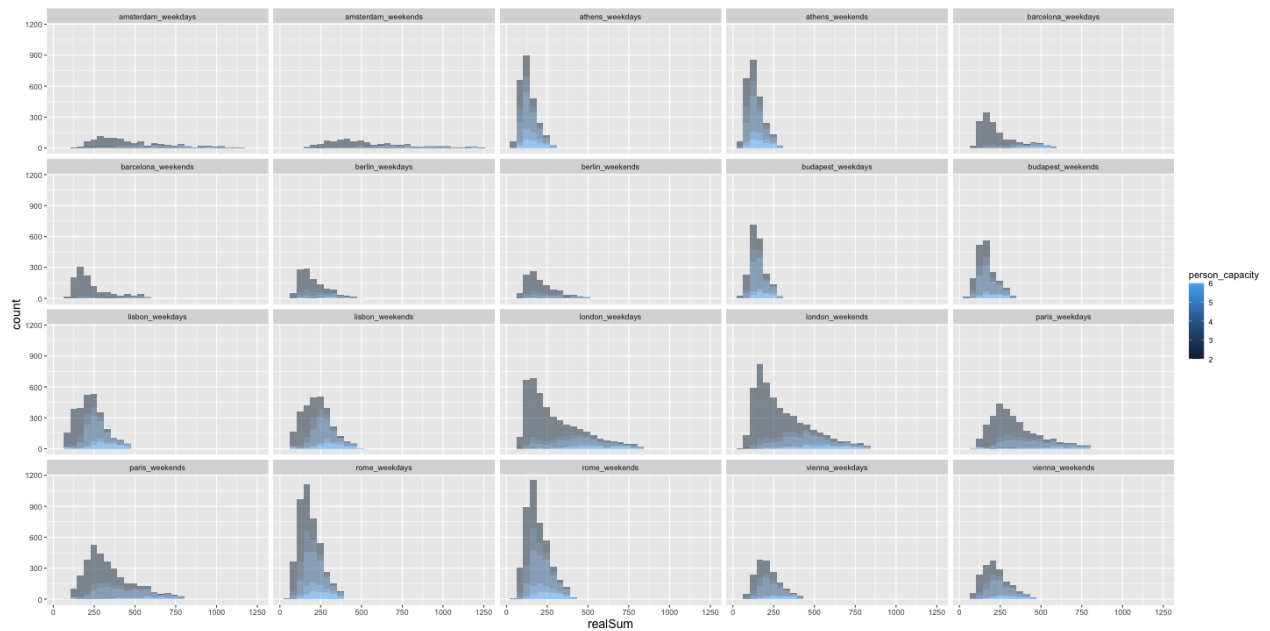
```
ggplot(my_data, aes(x = realSum, fill = person_capacity, group = person_capacity)) +
  geom_histogram(alpha = 0.5, nbins = 20) + theme(axis.title.x = element_text(size = 14),
axis.title.y = element_text(size = 14))
```



```
ggplot(my_data_filtered, aes(x = realSum, fill = person_capacity,
group = person_capacity)) + geom_histogram(alpha = 0.5, nbins = 20) +
  theme(axis.title.x = element_text(size = 14), axis.title.y = element_text(size = 14))
```




```
ggplot(my_data_filtered, aes(x = realSum, fill = person_capacity,
  group = person_capacity)) + geom_histogram(alpha = 0.5, nbins = 20) +
  theme(axis.title.x = element_text(size = 14), axis.title.y = element_text(size = 14)) +
  facet_wrap(~city_day)
```




The overall price is distributed similarly across the spectrum irrespective of person_capacity. But for some cities like london, london_weekdays, lisbon the price is higher with person capacity. So, person capacity along with city will be an important variable for determining price.

Real Sum Vs host_is_superuser


```
ggplot(my_data, aes(x = realSum, fill = host_is_superuser, group = host_is_superuser)) +  
  geom_histogram(alpha = 0.5, nbins = 20) + theme(axis.title.x = element_text(size = 14),  
  axis.title.y = element_text(size = 14))
```



```
ggplot(my_data_filtered, aes(x = realSum, fill = host_is_superuser,  
  group = host_is_superuser)) + geom_histogram(alpha = 0.5,  
  nbins = 20) + theme(axis.title.x = element_text(size = 14),  
  axis.title.y = element_text(size = 14))
```




```
ggplot(my_data_filtered, aes(x = realSum, fill = host_is_superhost,
  group = host_is_superhost)) + geom_histogram(alpha = 0.5,
  nbins = 20) + theme(axis.title.x = element_text(size = 14),
  axis.title.y = element_text(size = 14)) + facet_wrap(~city_day)
```




The prices are spread across all the spectrum irrespective of super_host or not.

Real Sum Vs multi


```
ggplot(my_data, aes(x = realSum, fill = multi, group = multi)) +
  geom_histogram(alpha = 0.5, nbins = 20) + theme(axis.title.x = element_text(size = 14),
  axis.title.y = element_text(size = 14))
```



```
ggplot(my_data_filtered, aes(x = realSum, fill = multi, group = multi)) +
  geom_histogram(alpha = 0.5, nbins = 20) + theme(axis.title.x = element_text(size = 14),
axis.title.y = element_text(size = 14))
```




```
ggplot(my_data_filtered, aes(x = realSum, fill = multi, group = multi)) +
  geom_histogram(alpha = 0.5, nbins = 20) + theme(axis.title.x = element_text(size = 14),
axis.title.y = element_text(size = 14)) + facet_wrap(~city_day)
```




The prices are similar irrespective of multi or not.

Real Sum Vs biz


```
ggplot(my_data, aes(x = realSum, fill = biz, group = biz)) +
  geom_histogram(alpha = 0.5, nbins = 20) + theme(axis.title.x = element_text(size = 14),
  axis.title.y = element_text(size = 14))
```



```
ggplot(my_data_filtered, aes(x = realSum, fill = biz, group = biz)) +
  geom_histogram(alpha = 0.5, nbins = 20) + theme(axis.title.x = element_text(size = 14),
  axis.title.y = element_text(size = 14))
```




```
ggplot(my_data_filtered, aes(x = realSum, fill = biz, group = biz)) +
  geom_histogram(alpha = 0.5, nbins = 20) + theme(axis.title.x = element_text(size = 14),
axis.title.y = element_text(size = 14)) + facet_wrap(~city_day)
```




The prices are similar irrespective of biz or not.

Real Sum vs Cleanliness

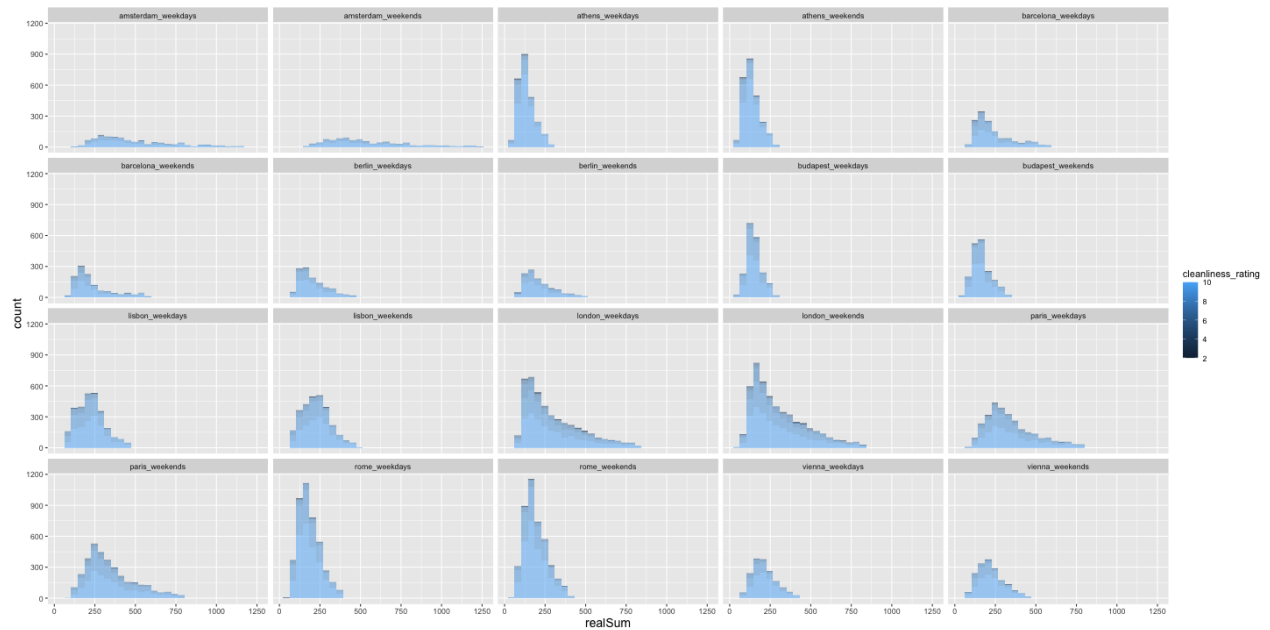
```
ggplot(my_data, aes(x = realSum, fill = cleanliness_rating, group = cleanliness_rating)) +
  geom_histogram(alpha = 0.5, nbins = 20) + theme(axis.title.x = element_text(size = 14),
axis.title.y = element_text(size = 14))
```



```
ggplot(my_data_filtered, aes(x = realSum, fill = cleanliness_rating,
  group = cleanliness_rating)) + geom_histogram(alpha = 0.5,
  nbins = 20) + theme(axis.title.x = element_text(size = 14),
  axis.title.y = element_text(size = 14))
```




```
ggplot(my_data_filtered, aes(x = realSum, fill = cleanliness_rating,
  group = cleanliness_rating)) + geom_histogram(alpha = 0.5,
  nbins = 20) + theme(axis.title.x = element_text(size = 14),
  axis.title.y = element_text(size = 14)) + facet_wrap(~city_day)
```

The cleanliness rating doesn't really have an effect on price

Scatterplot of Price vs Guest Satisfaction filtered by city

```
ggplot(my_data_filtered, aes(x = realSum, y = guest_satisfaction_overall,
  color = city_day)) + geom_point() + xlab("Price") + ylab("Guest Satisfaction Overall") +
  scale_color_discrete(name = "City-Day")
```




```
ggplot(my_data_filtered, aes(x = realSum, y = guest_satisfaction_overall,
  color = city_day)) + geom_point() + xlab("Price") + ylab("Guest Satisfaction Overall") +
  scale_color_discrete(name = "City-Day") + facet_wrap(~city_day)
```




The plot depicts that there is no correlation of price with guest satisfaction, good satisfaction rate is found across all the prices. In some cities like london, we can see a group of reviews with low guest satisfaction.

Real Sum Vs Bedroom Count

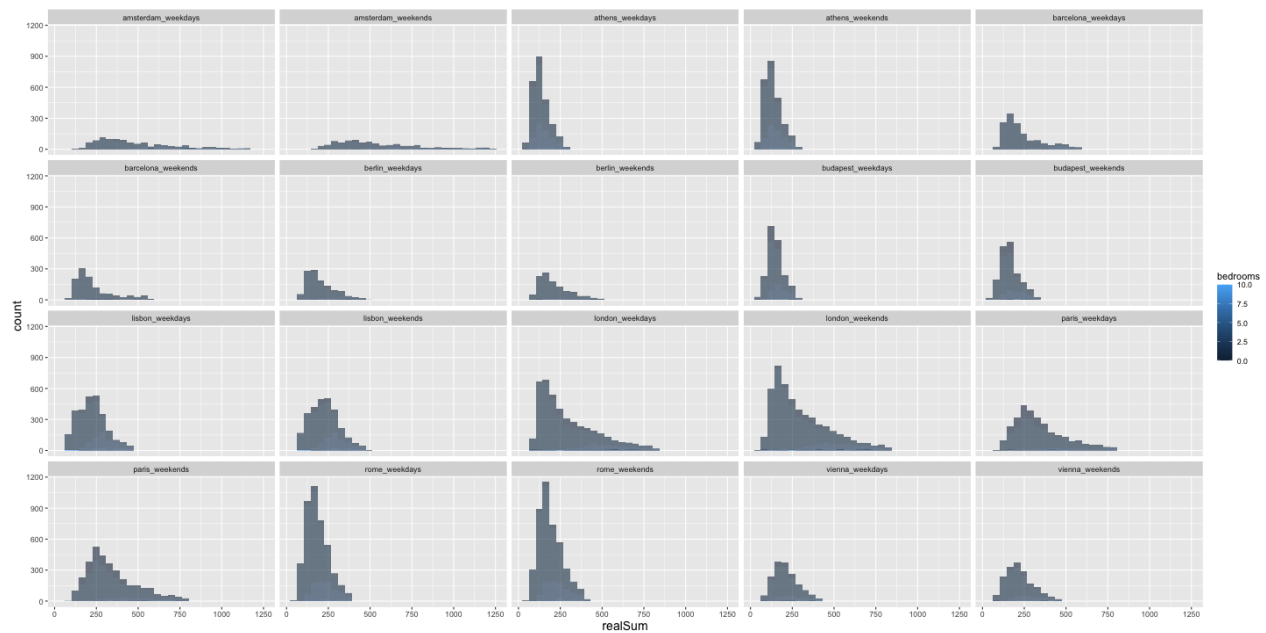
```
ggplot(my_data, aes(x = realSum, fill = bedrooms, group = bedrooms)) +
  geom_histogram(alpha = 0.6) + theme(axis.title.x = element_text(size = 14),
    axis.title.y = element_text(size = 14))
```




```
ggplot(my_data_filtered, aes(x = realSum, fill = bedrooms, group = bedrooms)) +
  geom_histogram(alpha = 0.6) + theme(axis.title.x = element_text(size = 14),
    axis.title.y = element_text(size = 14))
```



```
ggplot(my_data_filtered, aes(x = realSum, fill = bedrooms, group = bedrooms)) +
  geom_histogram(alpha = 0.6) + theme(axis.title.x = element_text(size = 14),
  axis.title.y = element_text(size = 14)) + facet_wrap(~city_day)
```



```
ggplot(my_data_filtered, aes(x = reorder(city_day, bedrooms,
FUN = median), y = bedrooms, fill = city_day)) + geom_boxplot() +
  coord_flip() + theme(legend.key.height = unit(0.5, "cm"),
  legend.key.size = unit(1, "lines"))
```




```
cor(as.numeric(factor(my_data$multi)), as.numeric(factor(my_data$biz)))
```

```
## [1] -0.4707248
```

Non Outlier Analysis

Boxplot of Price Vs City


```
ggplot(my_data_filtered, aes(x = reorder(city_day, realSum, FUN = median),
  y = realSum, fill = city_day)) + geom_boxplot() + coord_flip() +
  theme(legend.key.height = unit(0.5, "cm"), legend.key.size = unit(1,
    "lines"))
```



The highest prices in Europe are found in Amsterdam.

Density plot of Price vs Room type

```
ggplot(my_data_filtered, aes(x = realSum, group = room_type,
  fill = room_type, alpha = 0.2)) + geom_density()
```




The prices of entire home are high comparatively

Scatterplot of Prices in Rome w.r.t Latitude and Longitude during weekdays

```
tema <- theme(plot.title = element_text(size = 23, hjust = 0.5),
  axis.text.x = element_text(size = 19, face = "bold"), axis.text.y = element_text(size = 19,
    face = "bold"), axis.title.x = element_text(size = 19),
  axis.title.y = element_text(size = 19), legend.text = element_text(colour = "black",
    size = 19, face = "bold"), legend.background = element_rect(fill = "#F5FFFA",
    size = 0.5, linetype = "dashed", colour = "black"))

rome_data <- my_data_filtered %>%
  subset(city_day == "rome_weekdays")

ggplot(data = rome_data, mapping = aes(x = lat, y = lng)) + theme_minimal() +
  scale_fill_identity() + geom_point(mapping = aes(color = realSum),
  size = 3) + ggtitle("") + tema
```




This plot is within expectations of general trends, which suggests similar types of establishments (price and hospitality) tend to be in clusters.

Attraction Index in all Cities


```
# create a new column that groups the cities
my_data_filtered$city_group <- ifelse(my_data_filtered$city_day %in%
  c("amsterdam_weekdays", "amsterdam_weekends"), "amsterdam",
  ifelse(my_data_filtered$city_day %in% c("athens_weekdays",
    "athens_weekends"), "athens", ifelse(my_data_filtered$city_day %in%
    c("barcelona_weekdays", "barcelona_weekends"), "barcelona",
    ifelse(my_data_filtered$city_day %in% c("berlin_weekdays",
      "berlin_weekends"), "berlin", ifelse(my_data_filtered$city_day %in%
      c("budapest_weekdays", "budapest_weekends"), "budapest",
      ifelse(my_data_filtered$city_day %in% c("lisbon_weekdays",
        "lisbon_weekends"), "lisbon", ifelse(my_data_filtered$city_day %in%
        c("london_weekdays", "london_weekends"), "london",
        ifelse(my_data_filtered$city_day %in% c("paris_weekdays",
          "paris_weekends"), "paris", ifelse(my_data_filtered$city_day %in%
          c("rome_weekdays", "rome_weekends"), "rome",
          "vienna"))))))))

# plot the density plot with the new groupings
ggplot(my_data_filtered, aes(x = attr_index_norm, fill = city_group,
  alpha = 0.2)) + geom_density() + scale_fill_manual(values = c(amsterdam = "red",
  athens = "blue", barcelona = "green", berlin = "orange",
  budapest = "purple", lisbon = "black", london = "brown",
  paris = "grey", rome = "pink", vienna = "yellow")) + labs(fill = "City Group")
```



Real Sum vs Room Shared for all Cities

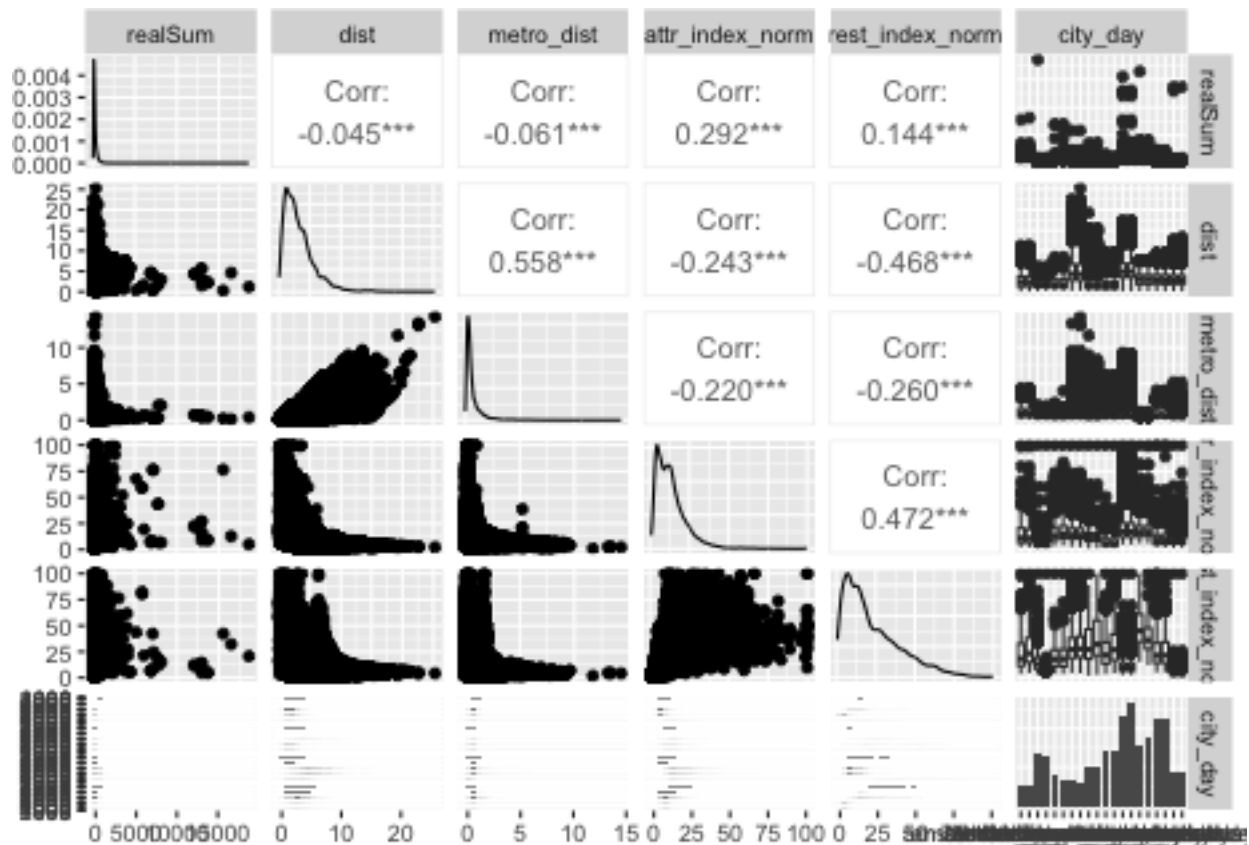
```
ggplot(my_data_filtered, aes(x = room_shared, y = realSum, color = city_group),
      alpha = 0.001) + geom_point() + scale_fill_manual(values = c(amsterdam = "red",
athens = "blue", barcelona = "green", berlin = "orange",
budapest = "purple", lisbon = "black", london = "brown",
paris = "grey", rome = "pink", vienna = "yellow")) + labs(fill = "City Group")
```


Different Model Selection and Training

Checking for correlations between different attributes

```
ggpairs(my_data[c("realSum", "dist", "metro_dist", "attr_index_norm",
  "rest_index_norm", "city_day")], cardinality = 20, cardinality_threshold = 999)
```




```
ggpairs(my_data_filtered[c("realSum", "dist", "metro_dist", "attr_index_norm",
"rest_index_norm", "city_day")], cardinality = 20, cardinality_threshold = 999)
```



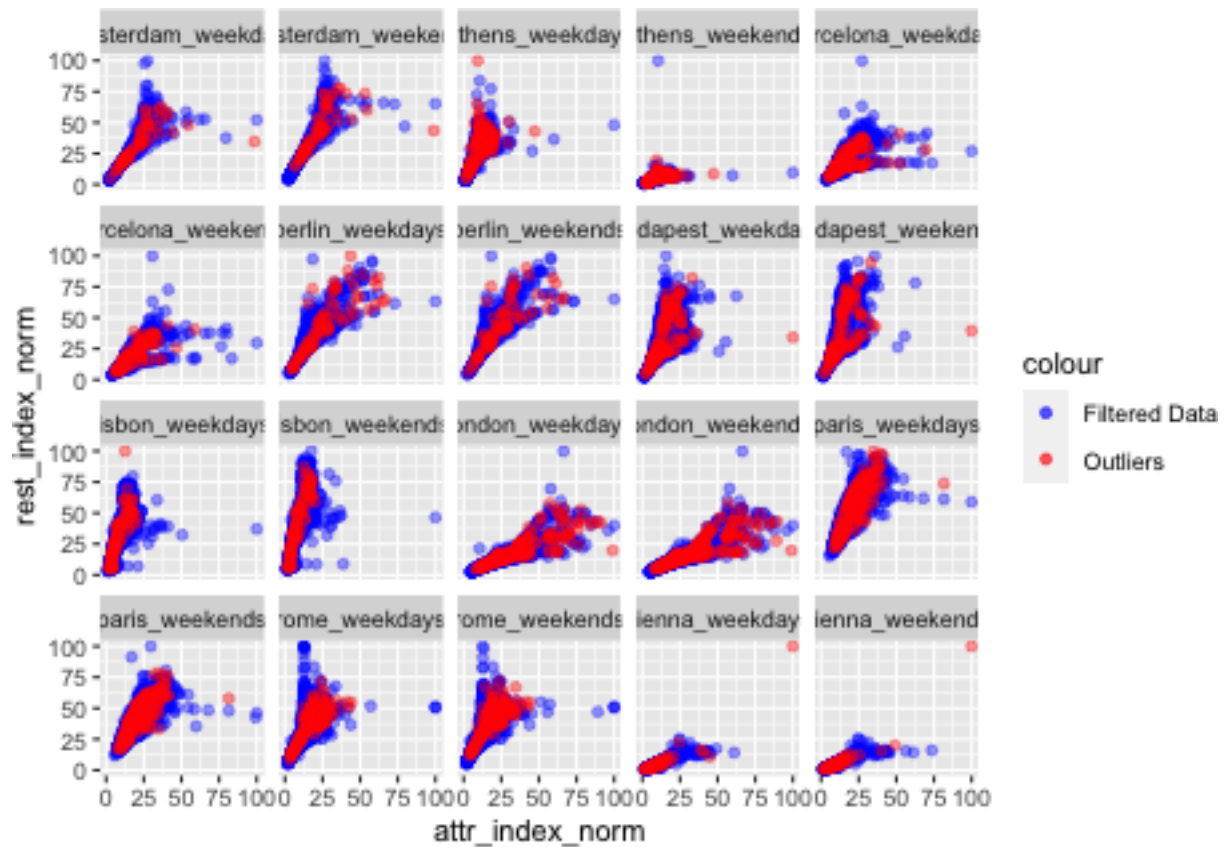
```
cor(my_data$attr_index, my_data$rest_index)
```

```
## [1] 0.4721427
```


```
ggplot() + geom_point(data = my_data_filtered, aes(x = attr_index_norm,
  y = rest_index_norm, color = "Filtered Data"), alpha = 0.4) +
  geom_point(data = my_outliers, aes(x = attr_index_norm, y = rest_index_norm,
    color = "Outliers"), alpha = 0.4) + scale_color_manual(values = c(`Filtered Data` = "blue",
    Outliers = "red"))
```




```
ggplot() + geom_point(data = my_data_filtered, aes(x = attr_index_norm,
  y = rest_index_norm, color = "Filtered Data"), alpha = 0.4) +
  geom_point(data = my_outliers, aes(x = attr_index_norm, y = rest_index_norm,
    color = "Outliers"), alpha = 0.4) + scale_color_manual(values = c(`Filtered Data` = "blue",
    Outliers = "red")) + facet_wrap(~city_day)
```



```
ggplot() + geom_point(data = my_data, aes(x = attr_index_norm,
  y = rest_index_norm, color = "Filtered Data"), alpha = 0.4) +
  scale_color_manual(values = c(`Filtered Data` = "blue"))
```



```
ggplot() + geom_point(data = my_data, aes(x = attr_index_norm,  
  y = rest_index_norm, color = "Filtered Data"), alpha = 0.4) +  
  scale_color_manual(values = c(`Filtered Data` = "blue")) +  
  facet_wrap(~city_day)
```



```
cor(my_data$attr_index_norm, my_data$rest_index_norm)
```

```
## [1] 0.4721427
```

Linear, Polynomial and Step Regression

MLR Seperated by City Day

```
temp_data <- subset(my_data_train, city_day == "amsterdam_weekends" |
  city_day == "amsterdam_weekdays")

M_0 <- lm(realSum ~ . - realSum - X, data = temp_data)
temp_data <- subset(my_data_train, city_day == "athens_weekdays" |
  city_day == "athens_weekends", )

M_1 <- lm(realSum ~ . - realSum - X - city_day, data = temp_data)
temp_data <- subset(my_data_train, city_day == "barcelona_weekdays" |
  city_day == "barcelona_weekends", )

M_2 <- lm(realSum ~ . - realSum - X - city_day, data = temp_data)
temp_data <- subset(my_data_train, city_day == "berlin_weekdays" |
  city_day == "berlin_weekends", )

M_3 <- lm(realSum ~ . - realSum - X - city_day, data = temp_data)
```

```

temp_data <- subset(my_data_train, city_day == "budapest_weekdays" |
  city_day == "budapest_weekends", )

M_4 <- lm(realSum ~ . - realSum - X - city_day, data = temp_data)
temp_data <- subset(my_data_train, city_day == "lisbon_weekdays" |
  city_day == "lisbon_weekends", )

M_5 <- lm(realSum ~ . - realSum - X - city_day, data = temp_data)
temp_data <- subset(my_data_train, city_day == "london_weekdays" |
  city_day == "london_weekends", )

M_6 <- lm(realSum ~ . - realSum - X - city_day, data = temp_data)
temp_data <- subset(my_data_train, city_day == "paris_weekdays" |
  city_day == "paris_weekends", )

M_7 <- lm(realSum ~ . - realSum - X - city_day, data = temp_data)
temp_data <- subset(my_data_train, city_day == "rome_weekdays" |
  city_day == "rome_weekends", )

M_8 <- lm(realSum ~ . - realSum - X - city_day, data = temp_data)
temp_data <- subset(my_data_train, city_day == "vienna_weekdays" |
  city_day == "vienna_weekends", )

M_9 <- lm(realSum ~ . - realSum - X - city_day, data = temp_data)

```

```

coefs <- tidy(M_0)
coefs[order(coefs$estimate, decreasing = TRUE), ]
coefs <- tidy(M_1)
coefs[order(coefs$estimate, decreasing = TRUE), ]
coefs <- tidy(M_2)
coefs[order(coefs$estimate, decreasing = TRUE), ]
coefs <- tidy(M_3)
coefs[order(coefs$estimate, decreasing = TRUE), ]
coefs <- tidy(M_4)
coefs[order(coefs$estimate, decreasing = TRUE), ]
coefs <- tidy(M_5)
coefs[order(coefs$estimate, decreasing = TRUE), ]
coefs <- tidy(M_6)
coefs[order(coefs$estimate, decreasing = TRUE), ]
coefs <- tidy(M_7)
coefs[order(coefs$estimate, decreasing = TRUE), ]
coefs <- tidy(M_8)
coefs[order(coefs$estimate, decreasing = TRUE), ]
coefs <- tidy(M_9)

```

```

coefs[order(coefs$estimate, decreasing = TRUE), ]

```

MLR

```

M1 <- lm(realSum ~ . - realSum - X, data = my_data_train)

```



```
summary(M1)
```

```
##
## Call:
## lm(formula = realSum ~ . - realSum - X, data = my_data_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -758.3   -84.0   -21.0    42.9  18422.4
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -4954.9076   3996.1824  -1.240  0.215017
## room_typePrivate room    -114.3655     4.2823  -26.707 < 2e-16 ***
## room_typeShared room   -204.1842    18.9348  -10.784 < 2e-16 ***
## person_capacity     23.9626     1.7645   13.581 < 2e-16 ***
## host_is_superhostTrue     1.0749     3.9344    0.273  0.784700
## multi           9.6004     4.1324    2.323  0.020173 *
## biz           33.2806     4.1885    7.946  1.99e-15 ***
## cleanliness_rating     5.0383     2.4153    2.086  0.036987 *
## guest_satisfaction_overall  0.7760     0.2615    2.968  0.002999 **
## bedrooms          86.0154     3.1888   26.974 < 2e-16 ***
## dist          -1.5330     1.2628   -1.214  0.224761
## metro_dist       -3.9967     2.5025   -1.597  0.110262
## attr_index_norm     6.3705     0.2946   21.627 < 2e-16 ***
## rest_index_norm    -0.1837     0.1774   -1.036  0.300215
## lng          -262.8909    40.1931   -6.541  6.20e-11 ***
## lat           123.2117    76.5228    1.610  0.107378
## city_dayamsterdam_weekends  67.9410    16.0017    4.246  2.18e-05 ***
## city_dayathens_weekdays  6315.0906   1388.5351    4.548  5.43e-06 ***
## city_dayathens_weekends  6303.5311   1388.6527    4.539  5.66e-06 ***
## city_daybarcelona_weekdays  411.7717    837.7909    0.491  0.623078
## city_daybarcelona_weekends  429.6529    837.8109    0.513  0.608075
## city_dayberlin_weekdays  1949.0424    342.1245    5.697  1.23e-08 ***
## city_dayberlin_weekends  1958.8844    342.0401    5.727  1.03e-08 ***
## city_daybudapest_weekdays  3902.3511    706.8806    5.521  3.40e-08 ***
## city_daybudapest_weekends  3929.6734    706.8440    5.559  2.73e-08 ***
## city_daylisbon_weekdays -2312.9309   1143.8977   -2.022  0.043186 *
## city_daylisbon_weekends -2304.0067   1143.8126   -2.014  0.043983 *
## city_daylondon_weekdays -1409.2997    206.1046   -6.838  8.17e-12 ***
## city_daylondon_weekends -1410.9328    206.1223   -6.845  7.76e-12 ***
## city_dayparis_weekdays  -403.0289    278.4819   -1.447  0.147840
## city_dayparis_weekends  -422.1389    278.6437   -1.515  0.129787
## city_dayrome_weekdays   2947.6852    881.3984    3.344  0.000826 ***
## city_dayrome_weekends   2952.5352    881.4297    3.350  0.000810 ***
## city_dayvienna_weekdays  3231.8185    582.7092    5.546  2.94e-08 ***
## city_dayvienna_weekends  3230.2680    582.7972    5.543  3.00e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 305.1 on 36159 degrees of freedom
## Multiple R-squared:  0.215, Adjusted R-squared:  0.2143
## F-statistic: 291.3 on 34 and 36159 DF, p-value: < 2.2e-16
```

```
# Create summary table with coefficients and p-values
table <- summary(M1)$coefficients[, c(1, 4)]
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(M1, my_data_train)
y_train_mean <- mean(my_data_train$realSum)
SST <- sum((my_data_train$realSum - y_train_mean)^2)
SSR <- sum((my_data_train$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_train$realSum)
p <- ncol(my_data_train)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_train$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.2150414
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.2146725
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 304.9175
```

The r^2 and adjusted r^2 values are too low for the Linear regression model to be considered a competent one in this case.

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(M1, my_data_test)
y_train_mean <- mean(my_data_test$realSum)
SST <- sum((my_data_test$realSum - y_train_mean)^2)
SSR <- sum((my_data_test$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_test$realSum)
p <- ncol(my_data_test)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_test$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.33744
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.3367131
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 233.2618
```

```
M1_step = step(M1, direction = "backward")
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(M1_step, my_data_train)
y_train_mean <- mean(my_data_train$realSum)
SST <- sum((my_data_train$realSum - y_train_mean)^2)
SSR <- sum((my_data_train$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_train$realSum)
p <- ncol(my_data_train)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_train$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.2149964
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.2146275
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 304.9262
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(M1_step, my_data_test)
y_train_mean <- mean(my_data_test$realSum)
SST <- sum((my_data_test$realSum - y_train_mean)^2)
SSR <- sum((my_data_test$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_test$realSum)
p <- ncol(my_data_test)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_test$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.3373441
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

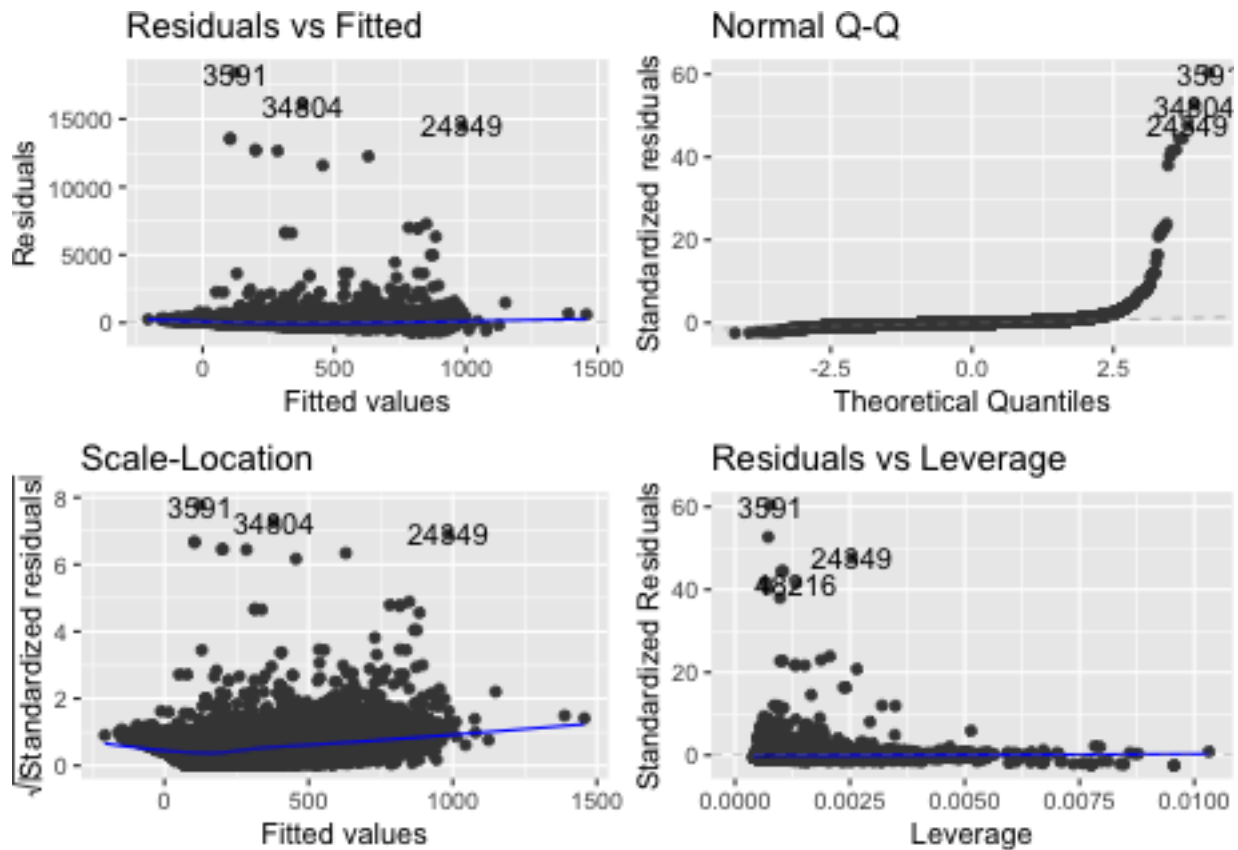
```
## Adjusted R-squared: 0.336617
```

```
cat("RMSE:", RMSE, "\n")
```

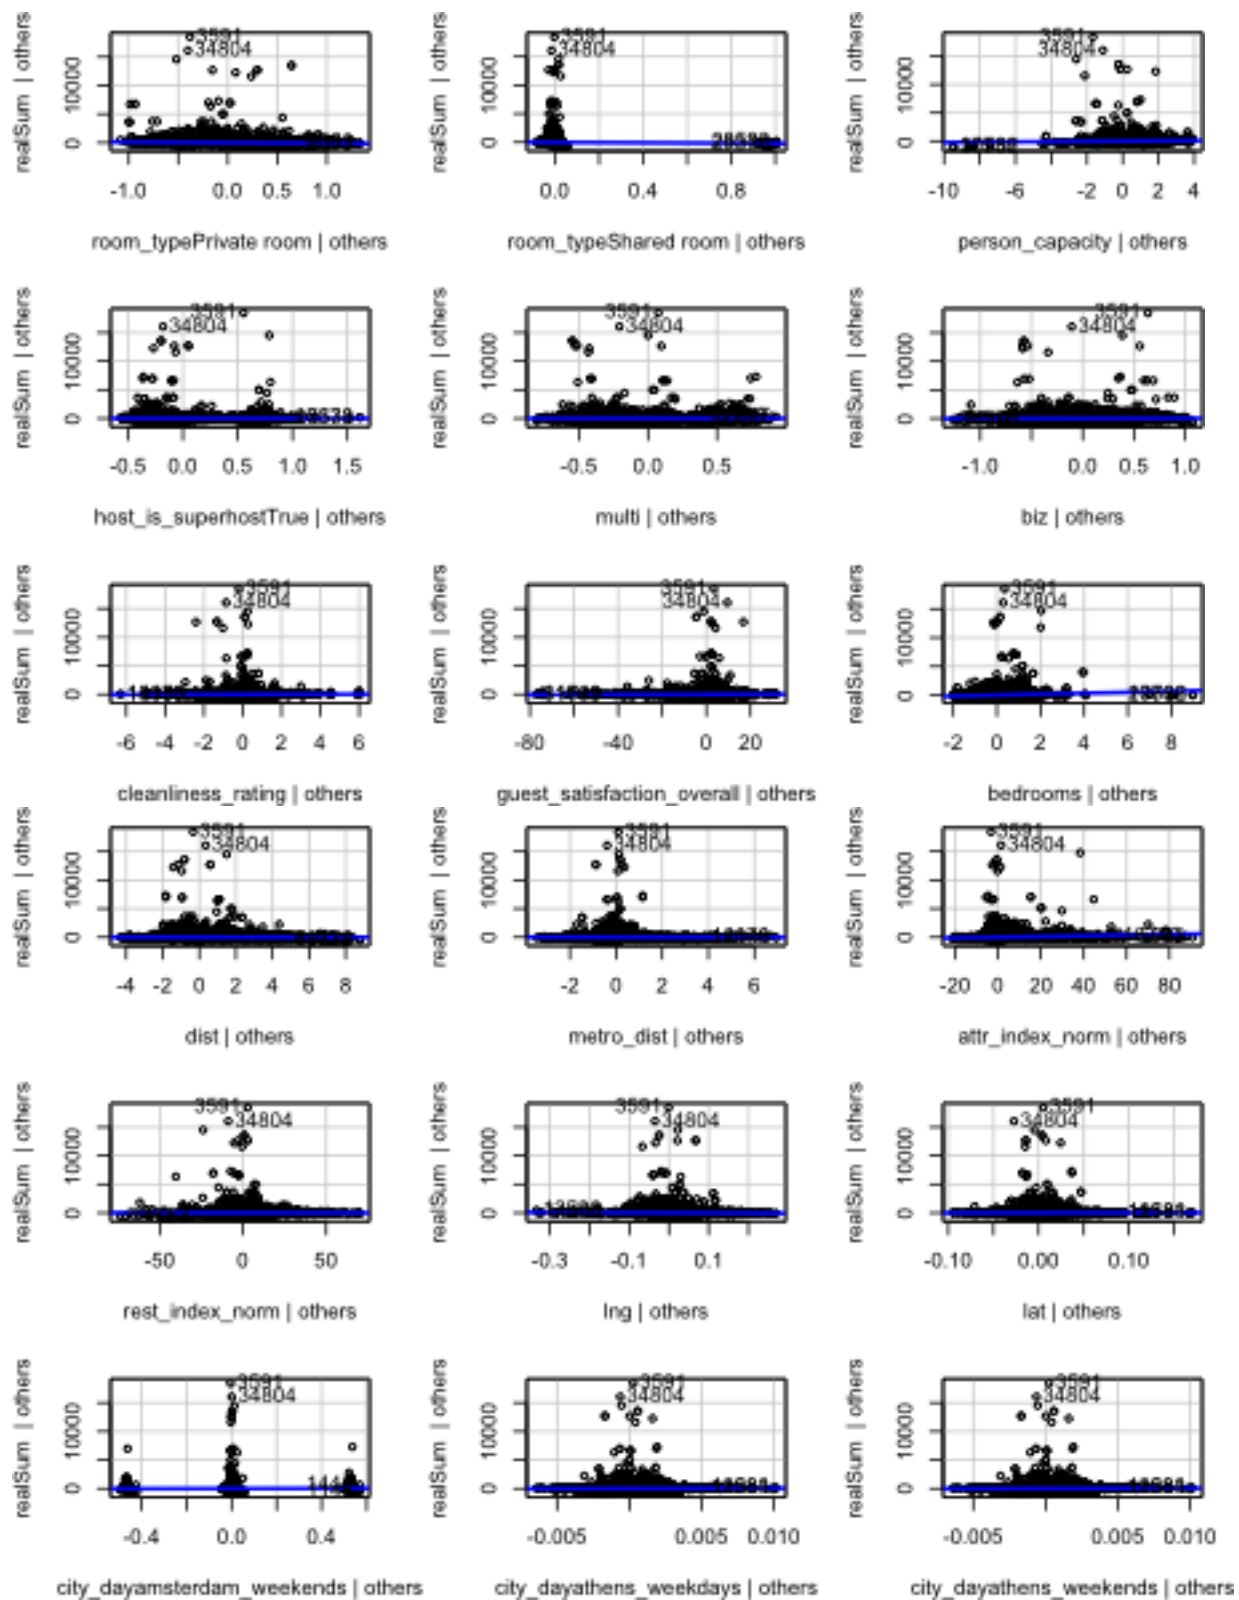
```
## RMSE: 233.2787
```

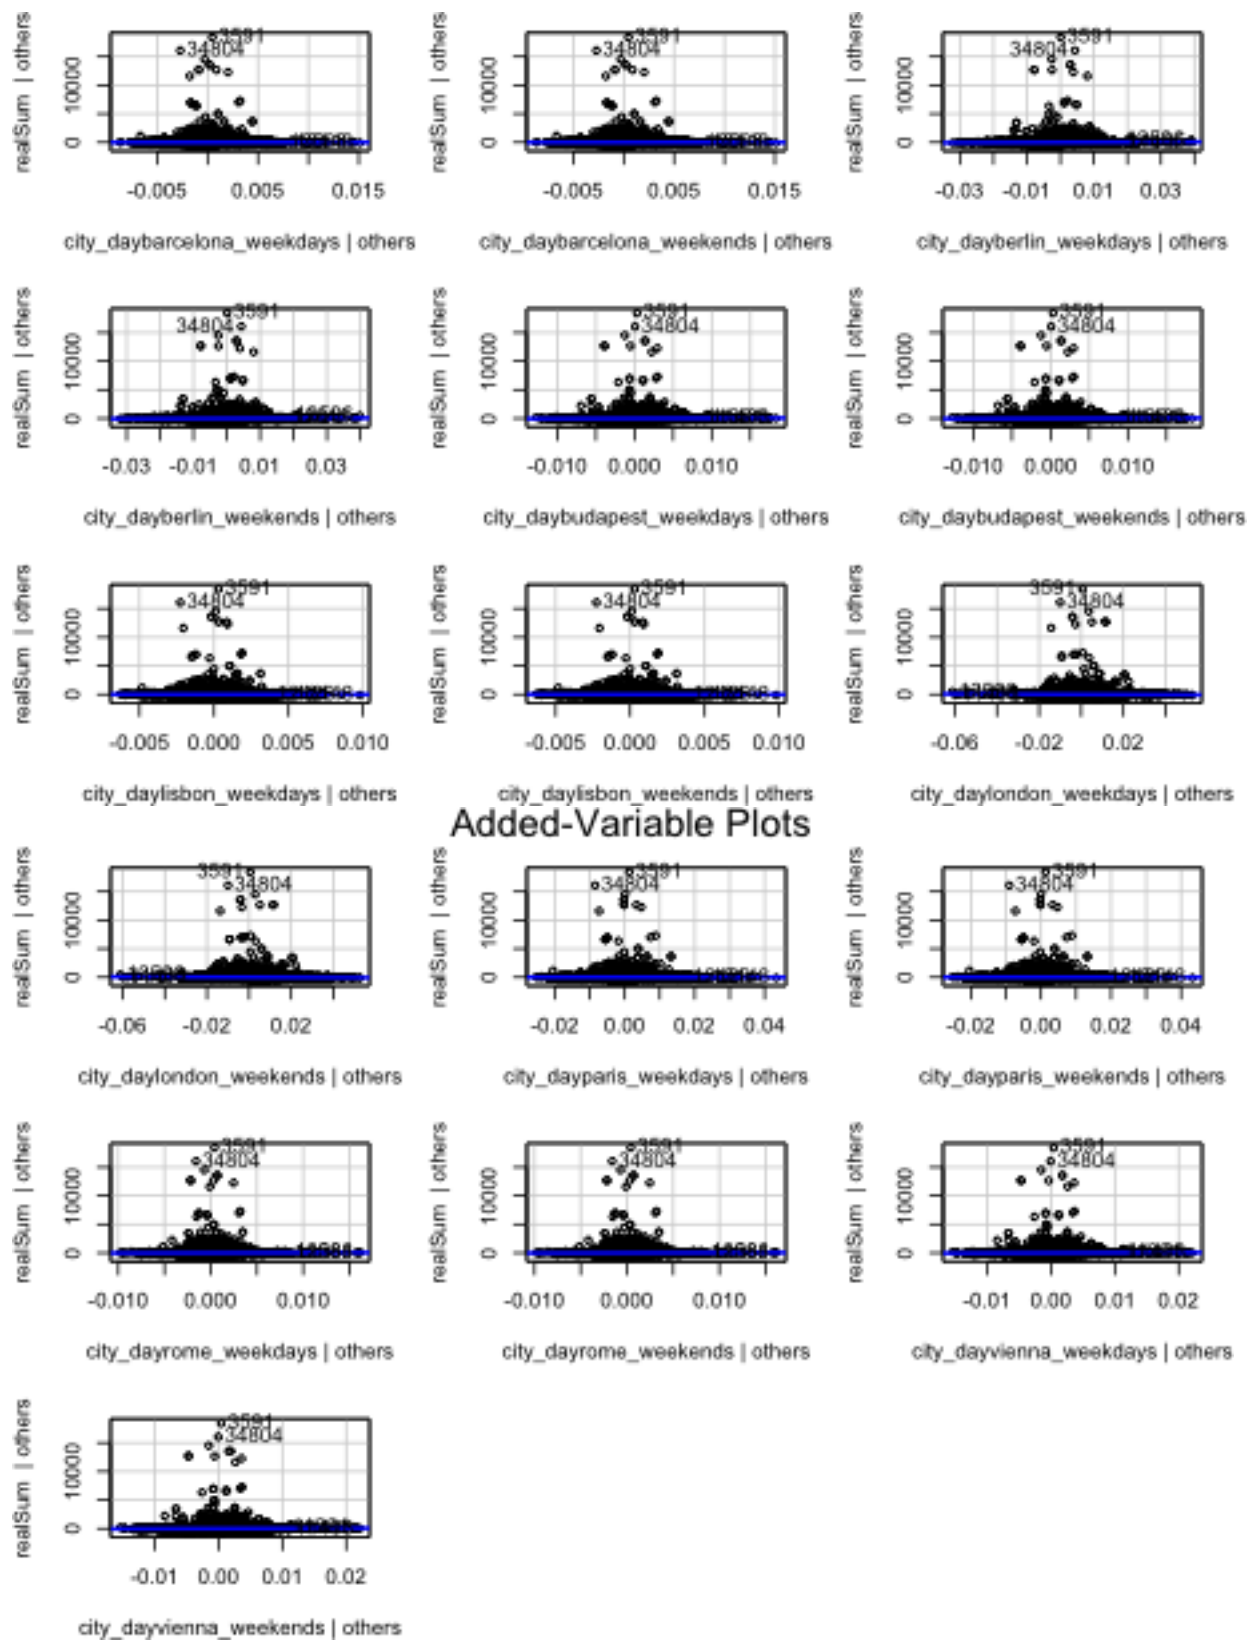
M1_step Diagnostics

```
library(car)  
autoplot(M1)
```



```
avPlots(M1)
```





```
my_data_train[34804, ]
```

```
##          X  realSum      room_type person_capacity host_is_superhost multi biz
## 49907 1736 637.6364 Entire home/apt                2          False    0    0
##      cleanliness_rating guest_satisfaction_overall bedrooms      dist
## 49907          10                93          1 0.9940386
##      metro_dist attr_index_norm rest_index_norm      lng      lat
## 49907 0.2025371      12.10842      6.748565 16.38568 48.2046
##          city_day
## 49907 vienna_weekdays
```

```
my_data_train[3591, ]
```

```
##          X  realSum      room_type person_capacity host_is_superhost multi biz
## 4917 183 156.3049 Entire home/apt                6          False    0    1
##      cleanliness_rating guest_satisfaction_overall bedrooms      dist metro_dist
## 4917          9                94          2 2.841432 0.04608507
##      attr_index_norm rest_index_norm      lng      lat      city_day
## 4917      2.366625      1.39056 23.72258 37.99906 athens_weekends
```

```
my_data_train[24349, ]
```

```
##          X  realSum      room_type person_capacity host_is_superhost multi biz
## 21461 1904 108.818 Private room                2          False    0    1
##      cleanliness_rating guest_satisfaction_overall bedrooms      dist
## 21461          8                80          1 3.216239
##      metro_dist attr_index_norm rest_index_norm      lng      lat
## 21461 0.3290175      3.115481      14.47335 -9.14292 38.74124
##          city_day
## 21461 lisbon_weekends
```

MLR with IVs

```
M1IV <- lm(realSum ~ room_type + host_is_superhost + multi +
  biz + city_day + person_capacity + cleanliness_rating + guest_satisfaction_overall +
  bedrooms + dist + metro_dist + attr_index_norm + rest_index_norm +
  lng + lat + metro_dist:dist + attr_index_norm:dist + attr_index_norm:metro_dist +
  rest_index_norm:dist + rest_index_norm:metro_dist + rest_index_norm:attr_index_norm,
  data = my_data_train)
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(M1IV, my_data_train)
y_train_mean <- mean(my_data_train$realSum)
SST <- sum((my_data_train$realSum - y_train_mean)^2)
SSR <- sum((my_data_train$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_train$realSum)
p <- ncol(my_data_train)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
```

```
RMSE = sqrt(mean((my_data_train$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values  
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.2159815
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.215613
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 304.7348
```

```
# Calculate R-squared and multiple R-squared  
y_train_pred <- predict(M1IV, my_data_test)  
y_train_mean <- mean(my_data_test$realSum)  
SST <- sum((my_data_test$realSum - y_train_mean)^2)  
SSR <- sum((my_data_test$realSum - y_train_pred)^2)  
R_squared <- 1 - SSR/SST  
n <- length(my_data_test$realSum)  
p <- ncol(my_data_test)  
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))  
RMSE = sqrt(mean((my_data_test$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values  
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.3379444
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.337218
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 233.173
```

```
M1stepIV = step(M1IV, direction = "backward")
```

```
# Calculate R-squared and multiple R-squared  
y_train_pred <- predict(M1stepIV, my_data_train)  
y_train_mean <- mean(my_data_train$realSum)  
SST <- sum((my_data_train$realSum - y_train_mean)^2)  
SSR <- sum((my_data_train$realSum - y_train_pred)^2)
```



```

R_squared <- 1 - SSR/SST
n <- length(my_data_train$realSum)
p <- ncol(my_data_train)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_train$realSum - y_train_pred)^2))

```

```

# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")

```

```
## R-squared: 0.2159435
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.215575
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 304.7422
```

```

# Calculate R-squared and multiple R-squared
y_train_pred <- predict(M1stepIV, my_data_test)
y_train_mean <- mean(my_data_test$realSum)
SST <- sum((my_data_test$realSum - y_train_mean)^2)
SSR <- sum((my_data_test$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_test$realSum)
p <- ncol(my_data_test)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_test$realSum - y_train_pred)^2))

```

```

# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")

```

```
## R-squared: 0.3380392
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.3373129
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 233.1563
```

Second Order Polynomial

```
poly2 <- lm(realSum ~ room_type + host_is_superhost + multi +
  biz + city_day + person_capacity + cleanliness_rating + guest_satisfaction_overall +
  bedrooms + poly(dist, 2) + poly(metro_dist, 2) + poly(attr_index_norm,
  2) + poly(rest_index_norm, 2) + lng + lat, data = my_data_train)
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(poly2, my_data_train)
y_train_mean <- mean(my_data_train$realSum)
SST <- sum((my_data_train$realSum - y_train_mean)^2)
SSR <- sum((my_data_train$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_train$realSum)
p <- ncol(my_data_train)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_train$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.2154699
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.2151012
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 304.8342
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(poly2, my_data_test)
y_train_mean <- mean(my_data_test$realSum)
SST <- sum((my_data_test$realSum - y_train_mean)^2)
SSR <- sum((my_data_test$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_test$realSum)
p <- ncol(my_data_test)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_test$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.3373538
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.3366268
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 233.277
```

```
poly2step = step(poly2, direction = "backward")
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(poly2step, my_data_train)
y_train_mean <- mean(my_data_train$realSum)
SST <- sum((my_data_train$realSum - y_train_mean)^2)
SSR <- sum((my_data_train$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_train$realSum)
p <- ncol(my_data_train)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_train$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.2154289
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.2150602
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 304.8422
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(poly2step, my_data_test)
y_train_mean <- mean(my_data_test$realSum)
SST <- sum((my_data_test$realSum - y_train_mean)^2)
SSR <- sum((my_data_test$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_test$realSum)
p <- ncol(my_data_test)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_test$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.3373204
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.3365933
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 233.2829
```

Second Order Polynomial with IVs

```
poly2IV <- lm(realSum ~ room_type + host_is_superhost + multi +  
  biz + city_day + person_capacity + cleanliness_rating + guest_satisfaction_overall +  
  bedrooms + poly(dist, 2) * poly(metro_dist, 2) * poly(attr_index_norm,  
  2) * poly(rest_index_norm, 2) + lng + lat, data = my_data_train)
```

```
# Calculate R-squared and multiple R-squared  
y_train_pred <- predict(poly2IV, my_data_train)  
y_train_mean <- mean(my_data_train$realSum)  
SST <- sum((my_data_train$realSum - y_train_mean)^2)  
SSR <- sum((my_data_train$realSum - y_train_pred)^2)  
R_squared <- 1 - SSR/SST  
n <- length(my_data_train$realSum)  
p <- ncol(my_data_train)  
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))  
RMSE = sqrt(mean((my_data_train$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values  
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.22214
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.2217744
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 303.5356
```

```
# Calculate R-squared and multiple R-squared  
y_train_pred <- predict(poly2IV, my_data_test)  
y_train_mean <- mean(my_data_test$realSum)  
SST <- sum((my_data_test$realSum - y_train_mean)^2)  
SSR <- sum((my_data_test$realSum - y_train_pred)^2)  
R_squared <- 1 - SSR/SST  
n <- length(my_data_test$realSum)  
p <- ncol(my_data_test)
```

```
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_test$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.334993
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.3342634
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 233.6922
```

```
poly2stepIV = step(poly2IV, direction = "backward")
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(poly2stepIV, my_data_train)
y_train_mean <- mean(my_data_train$realSum)
SST <- sum((my_data_train$realSum - y_train_mean)^2)
SSR <- sum((my_data_train$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_train$realSum)
p <- ncol(my_data_train)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_train$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.2221344
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.2217689
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 303.5367
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(poly2stepIV, my_data_test)
y_train_mean <- mean(my_data_test$realSum)
SST <- sum((my_data_test$realSum - y_train_mean)^2)
```

```
SSR <- sum((my_data_test$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_test$realSum)
p <- ncol(my_data_test)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_test$realSum - y_train_pred)^2))

# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.3350694
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.3343399
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 233.6788
```

Third Order Polynomial

```
poly3 <- lm(realSum ~ room_type + host_is_superhost + multi +
  biz + city_day + person_capacity + cleanliness_rating + guest_satisfaction_overall +
  bedrooms + poly(dist, 3) + poly(metro_dist, 3) + poly(attr_index_norm,
  3) + poly(rest_index_norm, 3) + lng + lat, data = my_data_train)
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(poly3, my_data_train)
y_train_mean <- mean(my_data_train$realSum)
SST <- sum((my_data_train$realSum - y_train_mean)^2)
SSR <- sum((my_data_train$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_train$realSum)
p <- ncol(my_data_train)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_train$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.2160624
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.215694
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 304.7191
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(poly3, my_data_test)
y_train_mean <- mean(my_data_test$realSum)
SST <- sum((my_data_test$realSum - y_train_mean)^2)
SSR <- sum((my_data_test$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_test$realSum)
p <- ncol(my_data_test)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_test$realSum - y_train_pred)^2))

# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.3375855
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.3368588
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 233.2362
```

```
poly3step = step(poly3, direction = "backward")
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(poly3step, my_data_train)
y_train_mean <- mean(my_data_train$realSum)
SST <- sum((my_data_train$realSum - y_train_mean)^2)
SSR <- sum((my_data_train$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_train$realSum)
p <- ncol(my_data_train)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_train$realSum - y_train_pred)^2))

# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.2160337
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.2156653
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 304.7247
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(poly3step, my_data_test)
y_train_mean <- mean(my_data_test$realSum)
SST <- sum((my_data_test$realSum - y_train_mean)^2)
SSR <- sum((my_data_test$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_test$realSum)
p <- ncol(my_data_test)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_test$realSum - y_train_pred)^2))

# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.3376519
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.3369253
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 233.2245
```

Third Order Polynomial with IVs

```
poly3IV <- lm(realSum ~ room_type + host_is_superhost + multi +
  biz + city_day + person_capacity + cleanliness_rating + guest_satisfaction_overall +
  bedrooms + poly(dist, 3) * poly(metro_dist, 3) * poly(attr_index_norm,
  3) * poly(rest_index_norm, 3) + lng + lat, data = my_data_train)
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(poly3IV, my_data_train)
y_train_mean <- mean(my_data_train$realSum)
SST <- sum((my_data_train$realSum - y_train_mean)^2)
SSR <- sum((my_data_train$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_train$realSum)
p <- ncol(my_data_train)
```



```
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_train$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.2330663
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.2327059
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 301.3962
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(poly3IV, my_data_test)
y_train_mean <- mean(my_data_test$realSum)
SST <- sum((my_data_test$realSum - y_train_mean)^2)
SSR <- sum((my_data_test$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_test$realSum)
p <- ncol(my_data_test)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_test$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.1901115
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.189223
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 257.8955
```

```
poly3stepIV = step(poly3IV, direction = "backward")
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(poly3stepIV, my_data_train)
y_train_mean <- mean(my_data_train$realSum)
SST <- sum((my_data_train$realSum - y_train_mean)^2)
```

```
SSR <- sum((my_data_train$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_train$realSum)
p <- ncol(my_data_train)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_train$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.2285384
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.2281759
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 302.2846
```

```
# Calculate R-squared and multiple R-squared
y_train_pred <- predict(poly3stepIV, my_data_test)
y_train_mean <- mean(my_data_test$realSum)
SST <- sum((my_data_test$realSum - y_train_mean)^2)
SSR <- sum((my_data_test$realSum - y_train_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(my_data_test$realSum)
p <- ncol(my_data_test)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
RMSE = sqrt(mean((my_data_test$realSum - y_train_pred)^2))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.3281442
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.327407
```

```
cat("RMSE:", RMSE, "\n")
```

```
## RMSE: 234.8925
```

Lasso Regression

```


# Prepare the predictors and response variable
x_train <- model.matrix(realSum ~ room_type + person_capacity +
  host_is_superhost + multi + biz + cleanliness_rating + guest_satisfaction_overall +
  bedrooms + dist * metro_dist * attr_index_norm * rest_index_norm +
  lng + lat + city_day, data = my_data_train)[, -1]
y_train <- my_data_train$realSum
y_test <- my_data_test$realSum

# Fit a Lasso regression model
lasso_model <- glmnet(x_train, y_train, alpha = 1)

# Select the best lambda value using cross-validation
cv_model <- cv.glmnet(x_train, y_train, alpha = 1, nfolds = 5)

# Plot the cross-validation results
plot(cv_model)

```



```

# Select the lambda value that minimizes the mean
# cross-validation error
best_lambda <- cv_model$lambda.min

# Fit a Lasso regression model with the selected lambda
# value
lasso_model_best <- glmnet(x_train, y_train, alpha = 1, lambda = best_lambda)

# Calculate R-squared and multiple R-squared
y_train_pred <- predict(lasso_model_best, newx = x_train)
y_train_mean <- mean(y_train)
SST <- sum((y_train - y_train_mean)^2)

```

```
SSR <- sum((y_train - y_train_pred)^2)
R_squared <- 1 - SSR/SST
multiple_R_squared <- cor(y_train_pred, y_train)^2
n <- length(y_train)
p <- ncol(x_train)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.2157005
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.2147241
```

```
# Evaluate the model performance
rmse <- sqrt(mean((my_data_test$realSum - y_train_pred)^2))
cat("RMSE on train set:", rmse, "\n")
```

```
## RMSE on train set: 325.8081
```

```
x_test <- model.matrix(realSum ~ room_type + person_capacity +
  host_is_superhost + multi + biz + cleanliness_rating + guest_satisfaction_overall +
  bedrooms + dist * metro_dist * attr_index_norm * rest_index_norm +
  lng + lat + city_day, data = my_data_test)[, -1]
```

```
# Calculate R-squared and multiple R-squared
y_test_pred <- predict(lasso_model_best, newx = x_test)
y_test_mean <- mean(y_test)
SST <- sum((y_test - y_test_mean)^2)
SSR <- sum((y_test - y_test_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(y_test)
p <- ncol(x_test)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))
```

```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.3372477
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.3353195
```

```
# Evaluate the model performance
rmse <- sqrt(mean((my_data_test$realSum - y_test_pred)^2))
cat("RMSE on test set:", rmse, "\n")
```


```
## RMSE on test set: 233.2957
```

```
# Prepare the predictors and response variable
x_train <- model.matrix(realSum ~ room_type + person_capacity +
  host_is_superhost + multi + biz + cleanliness_rating + guest_satisfaction_overall +
  bedrooms + poly(dist, 2) * poly(metro_dist, 2) * poly(attr_index_norm,
    2) * poly(rest_index_norm, 2) + lng + lat + city_day, data = my_data_train)[,
  -1]
y_train <- my_data_train$realSum
y_test <- my_data_test$realSum

# Fit a Lasso regression model
lasso_model <- glmnet(x_train, y_train, alpha = 1)

# Select the best lambda value using cross-validation
cv_model <- cv.glmnet(x_train, y_train, alpha = 1, nfolds = 5)

# Plot the cross-validation results
plot(cv_model)
```



```
# Select the lambda value that minimizes the mean
# cross-validation error
best_lambda <- cv_model$lambda.min

# Fit a Lasso regression model with the selected lambda
```

```

# value
lasso_model_best <- glmnet(x_train, y_train, alpha = 1, lambda = best_lambda)

# Calculate R-squared and multiple R-squared
y_train_pred <- predict(lasso_model_best, newx = x_train)
y_train_mean <- mean(y_train)
SST <- sum((y_train - y_train_mean)^2)
SSR <- sum((y_train - y_train_pred)^2)
R_squared <- 1 - SSR/SST
multiple_R_squared <- cor(y_train_pred, y_train)^2
n <- length(y_train)
p <- ncol(x_train)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))

# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")

```

```
## R-squared: 0.2176971
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.2153122
```

```

# Evaluate the model performance
rmse <- sqrt(mean((my_data_test$realSum - y_train_pred)^2))
cat("RMSE on train set:", rmse, "\n")

```

```
## RMSE on train set: 324.7053
```

```

x_test <- model.matrix(realSum ~ room_type + person_capacity +
  host_is_superhost + multi + biz + cleanliness_rating + guest_satisfaction_overall +
  bedrooms + poly(dist, 2) * poly(metro_dist, 2) * poly(attr_index_norm,
  2) * poly(rest_index_norm, 2) + lng + lat + city_day, data = my_data_test)[,
  -1]

```

```

# Calculate R-squared and multiple R-squared
y_test_pred <- predict(lasso_model_best, newx = x_test)
y_test_mean <- mean(y_test)
SST <- sum((y_test - y_test_mean)^2)
SSR <- sum((y_test - y_test_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(y_test)
p <- ncol(x_test)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))

# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")

```

```
## R-squared: 0.2821967
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.2770702
```

```
# Evaluate the model performance
```

```
rmse <- sqrt(mean((my_data_test$realSum - y_test_pred)^2))  
cat("RMSE on test set:", rmse, "\n")
```

```
## RMSE on test set: 242.7917
```

```
# Prepare the predictors and response variable
```

```
x_train <- model.matrix(realSum ~ room_type + person_capacity +  
  host_is_superhost + multi + biz + cleanliness_rating + guest_satisfaction_overall +  
  bedrooms + poly(dist, 3) * poly(metro_dist, 3) * poly(attr_index_norm,  
  3) * poly(rest_index_norm, 3) + lng + lat + city_day, data = my_data_train)[,  
  -1]  
y_train <- my_data_train$realSum  
y_test <- my_data_test$realSum
```

```
# Fit a Lasso regression model
```


```
lasso_model <- glmnet(x_train, y_train, alpha = 1)
```

```
# Select the best lambda value using cross-validation
```

```
cv_model <- cv.glmnet(x_train, y_train, alpha = 1, nfolds = 5)
```

```
# Plot the cross-validation results
```

```
plot(cv_model)
```



```

# Select the lambda value that minimizes the mean
# cross-validation error
best_lambda <- cv_model$lambda.min

# Fit a Lasso regression model with the selected lambda
# value
lasso_model_best <- glmnet(x_train, y_train, alpha = 1, lambda = best_lambda)

# Calculate R-squared and multiple R-squared
y_train_pred <- predict(lasso_model_best, newx = x_train)
y_train_mean <- mean(y_train)
SST <- sum((y_train - y_train_mean)^2)
SSR <- sum((y_train - y_train_pred)^2)
R_squared <- 1 - SSR/SST
multiple_R_squared <- cor(y_train_pred, y_train)^2
n <- length(y_train)
p <- ncol(x_train)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))

# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")

```

```
## R-squared: 0.2188426
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.2126426
```

```

# Evaluate the model performance
rmse <- sqrt(mean((my_data_test$realSum - y_train_pred)^2))
cat("RMSE on train set:", rmse, "\n")

```

```
## RMSE on train set: 324.4427
```

```

x_test <- model.matrix(realSum ~ room_type + person_capacity +
  host_is_superhost + multi + biz + cleanliness_rating + guest_satisfaction_overall +
  bedrooms + poly(dist, 3) * poly(metro_dist, 3) * poly(attr_index_norm,
  3) * poly(rest_index_norm, 3) + lng + lat + city_day, data = my_data_test)[,
  -1]

# Calculate R-squared and multiple R-squared
y_test_pred <- predict(lasso_model_best, newx = x_test)
y_test_mean <- mean(y_test)
SST <- sum((y_test - y_test_mean)^2)
SSR <- sum((y_test - y_test_pred)^2)
R_squared <- 1 - SSR/SST
n <- length(y_test)
p <- ncol(x_test)
adj_R_squared <- 1 - (SSR/(n - p - 1))/(SST/(n - 1))

```



```
# Print the R-squared and multiple R-squared values
cat("R-squared:", R_squared, "\n")
```

```
## R-squared: 0.2587531
```

```
cat("Adjusted R-squared:", adj_R_squared, "\n")
```

```
## Adjusted R-squared: 0.2448794
```

```
# Evaluate the model performance
rmse <- sqrt(mean((my_data_test$realSum - y_test_pred)^2))
cat("RMSE on test set:", rmse, "\n")
```

```
## RMSE on test set: 246.7246
```

Even step regression is not good because of extremely low value of R^2 even in polynomial model of power 2 and 3.

Random Forest Regression

```
rf_model <- randomForest(realSum ~ room_type + host_is_superhost +
  multi + biz + city_day + person_capacity + cleanliness_rating +
  guest_satisfaction_overall + bedrooms + dist + metro_dist +
  attr_index_norm + rest_index_norm + lng + lat, data = my_data_train)
```

```
predictions <- predict(rf_model, my_data_train)
```

```
rmse <- sqrt(mean((my_data_train$realSum - predictions)^2))
rmse
```

```
## [1] 120.6731
```

```
r_squared <- 1 - (sum((my_data_train$realSum - predictions)^2)/sum((my_data_train$realSum -
  mean(my_data_train$realSum))^2))
r_squared
```

```
## [1] 0.8770571
```

```
adj_R_squared <- 1 - ((1 - r_squared) * (n - 1)/(n - p - 1))
adj_R_squared
```

```
## [1] 0.874756
```

```
predictions <- predict(rf_model, my_data_test)
```

```
rmse <- sqrt(mean((my_data_test$realSum - predictions)^2))
rmse
```

```
## [1] 140.1627
```

```
r_squared <- 1 - (sum((my_data_test$realSum - predictions)^2)/sum((my_data_test$realSum -
  mean(my_data_test$realSum))^2))
r_squared
```

```
## [1] 0.7607771
```

```
adj_R_squared <- 1 - ((1 - r_squared) * (n - 1)/(n - p - 1))
adj_R_squared
```

```
## [1] 0.7562996
```

```
importance(rf_model)
```

```
##                               IncNodePurity
## room_type                      128820124
## host_is_superhost              29206740
## multi                          23032640
## biz                           28483381
## city_day                      119817452
## person_capacity               181945316
## cleanliness_rating            84896646
## guest_satisfaction_overall    153697848
## bedrooms                     339811264
## dist                          403842463
## metro_dist                   260251950
## attr_index_norm              668769759
## rest_index_norm              397578038
## lng                          489702071
## lat                          605397008
```

```
varImpPlot(rf_model)
```

rf_model

